

Evolutionary Computation Meets Machine Learning for Combinatorial Optimisation

Yi Mei¹ and Günther R. Raidl²

1 Victoria University of Wellington, New Zealand

2 TU Wien, Vienna, Austria

yi.mei@ecs.vuw.ac.nz
raidl@ac.tuwien.ac.at



<http://gecco-2025.sigevo.org/>

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

GECCO '25 Companion, July 14 - 18, 2025, Málaga, Spain
© 2025 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-1464-1/25/07... \$15.00
<https://doi.org/10.1145/3712255.3716512>



Outline



- Introduction and Background
- Evolutionary Computation to Learn Combinatorial Optimisation Heuristics
- Machine Learning to Aid (Meta-)Heuristics
- Challenges and Future Directions

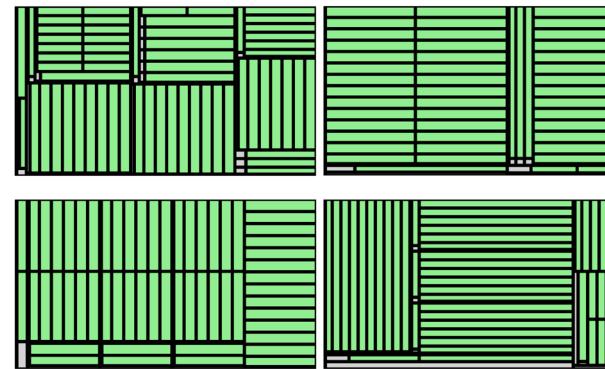
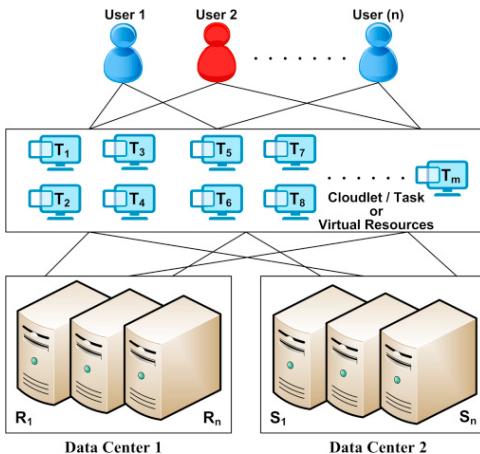
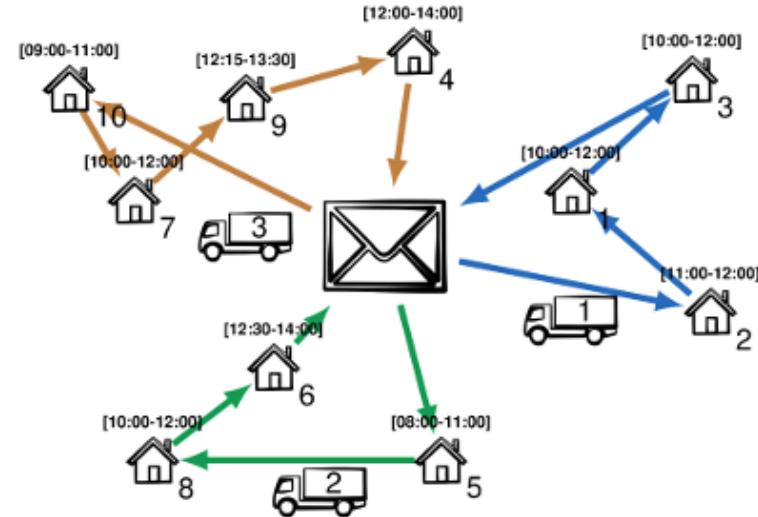


Introduction and Background

Combinatorial Optimisation

Many real-world applications

- Supply chain/logistics
- Cloud computing
- Manufacturing, scheduling
- Cutting and packing
- Bioinformatics
- ...
- *NP-hard, many local optima...*

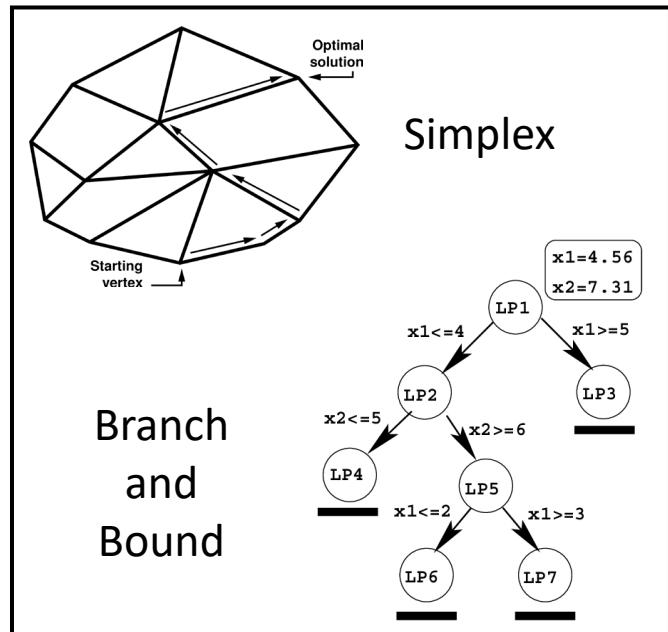


Combinatorial Optimisation Methods

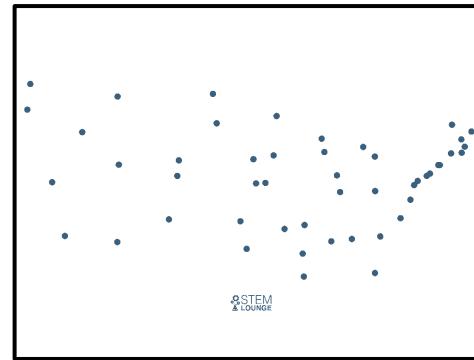
Exact Methods

Constructive
heuristics

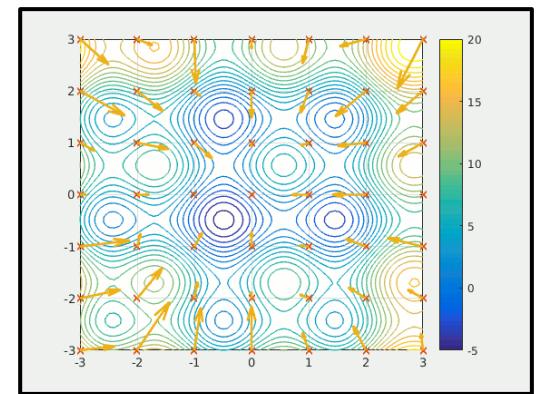
Meta-heuristics



Integer Linear
Programming



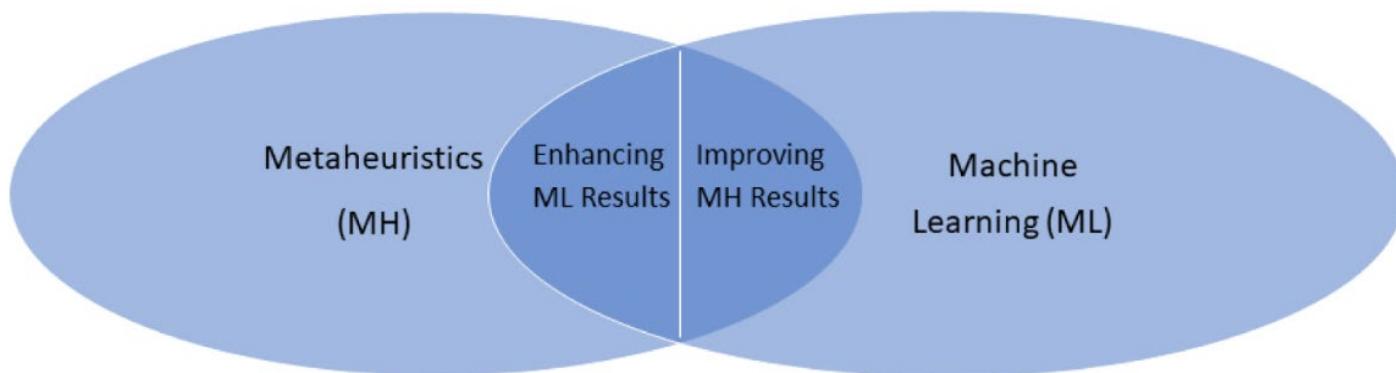
Nearest neighbour
heuristic for TSP



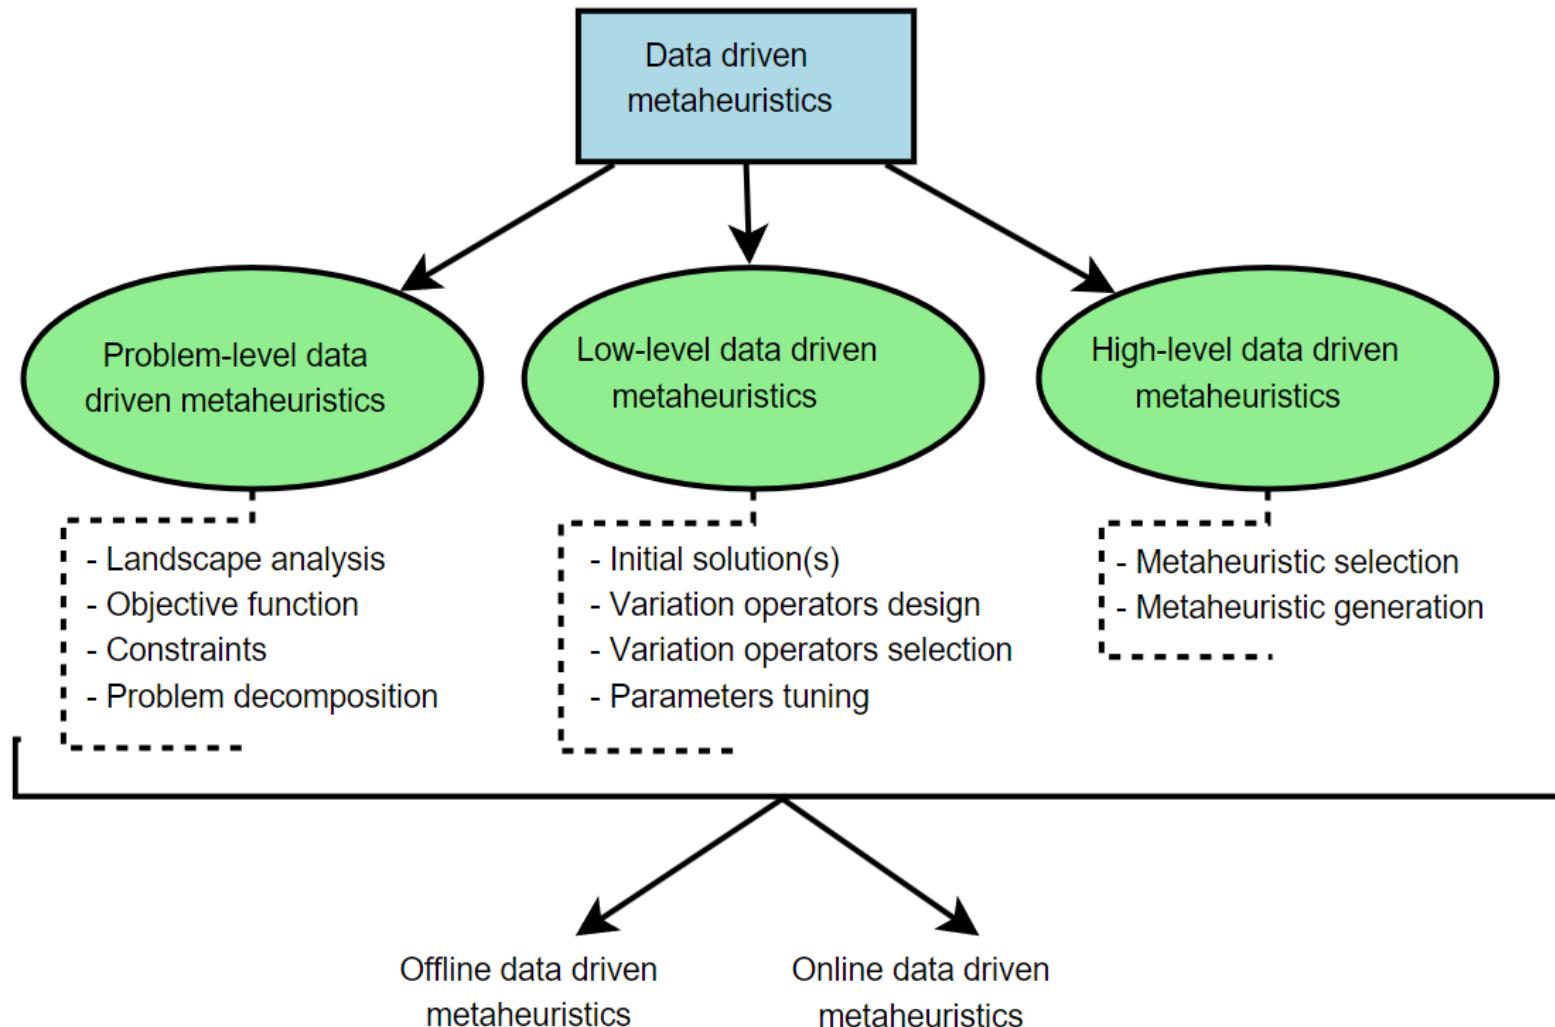
Particle Swarm
Optimization

Evolutionary Computation, Metaheuristics, and Learning

- AI/machine learning (ML) boom also hit the area of combinatorial optimization (COP)
- Evolutionary Computation and other (meta-)heuristics relate to machine learning in many different ways



Learning to (Better) Optimize



Survey Articles

- M. Karimi-Mamaghan, M. Mohammadi, P. Meyer, A. M. Karimi-Mamaghan, E.-G. Talbi (2021): Machine Learning at the Service of Metaheuristics for solving COPs: A State-of-the Art, EJOR
- E. G.Talbi (2021): Machine Learning into Metaheuristics: A Survey and Taxonomy, ACM Computing Surveys
- Y. Bengio, A. Lodi, A. Prouvost (2021): Machine Learning for Combinatorial Optimization: A Methodological Tour d'Horizon; EJOR
- N. Mazyavkina, S. Sviridov, S. Ivanov, E. Burnaev (2021): Reinforcement Learning for Combinatorial Optimization: A Survey; COR
- C. Zang, Y. Wu, Y. Ma, W. Song, Z. Le (2023): A Review on Learning to Solve Combinatorial Optimization Problems in Manufacturing, IET Collaborative Intelligent Manufacturing

Some Classical Metaheuristics Involving Learning

Basic idea of learning in (meta-)heuristics not new:

- Evolution Strategies
- Reactive Tabu Search
- Guided Local Search
- Variable Neighborhood Search,
Adaptive Large Neighborhood Search
 - self-adaptive selection of neighborhood structures/operators
- Hyper-Heuristics
- Ant Colony Optimization



Q-Learning (Watkins and Dayan, 1992)

- for discrete state and action spaces, off-policy RL
- Goal: max. weighted sum of expected rewards over all future steps
- Table with Q-value $Q(s, a)$ for each state/action pair (s, a) : approximates expected weighted sum of rewards to receive from state s onward when performing action a
- Update of initially random $Q(s, a)$ values by

$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \right)}_{\text{temporal difference}}$$

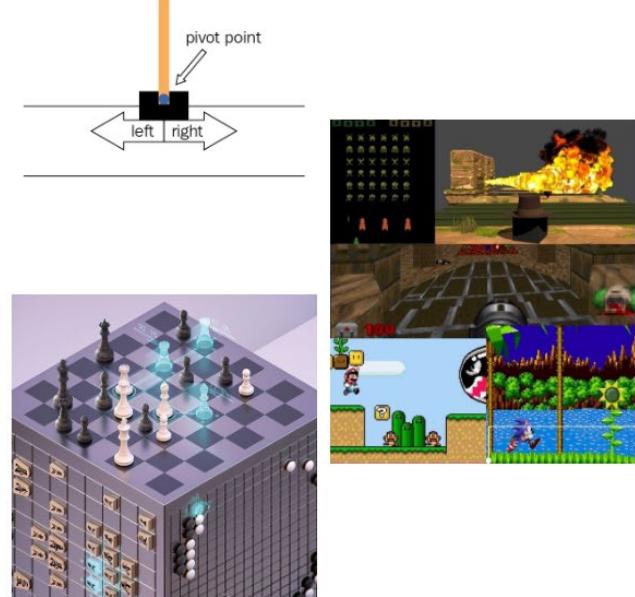
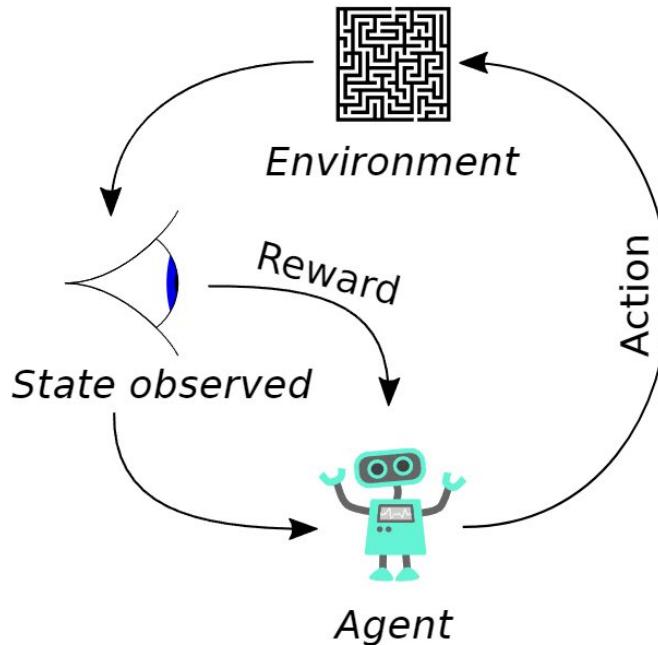
new value (temporal difference target)

- Compare with Ant Colony Optimization!
- Frequently too many states
 - approximate $Q(s, a)$ by a (deep) neural network
 - Deep Q Network (DQN, Mnih et al., 2013)



Reinforcement Learning

- A sub-discipline of machine learning
- Learn to behave well in a given **environment**
 - usually considered a Markov Decision Process (MDP)
- Framework:



Machine Learning for CO

- Learning to **construct** (initial) solutions
- Learning to **search** in the solution space
- **Supervised** learning
 - Collect **data** on problem instances + optimal solutions/steps
 - Learn to **mimic** the optimal solutions/steps
- **Reinforcement** learning
 - Given **training problem instances**, no solution needed
 - Learn a **policy** to make decisions
 - **Reward/Quality** indicators
 - How good the policy is
 - Based on solutions obtained by the policy



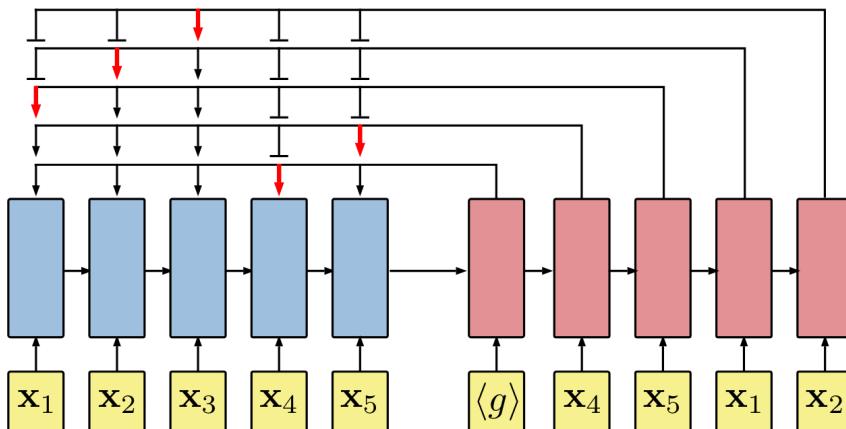
Learning to Construct

- End-to-end learning
 - Input: problem instance
 - Output: a solution
- Usually outputs the final solution, but may also be used as initial solution followed by a search process
- Key design issues:
 - Encoding of the problem instance (input)
 - Encoding of the solution (output)
 - Model architecture
 - Learning method



Example: RNN for TSP

- **Problem:** Find the shortest cycle to visit all the nodes in a graph
- **Design:**
 - **Input:** nodes as a sequence
 - **Output:** solution as a sequence
 - **Model:** recurrent neural network



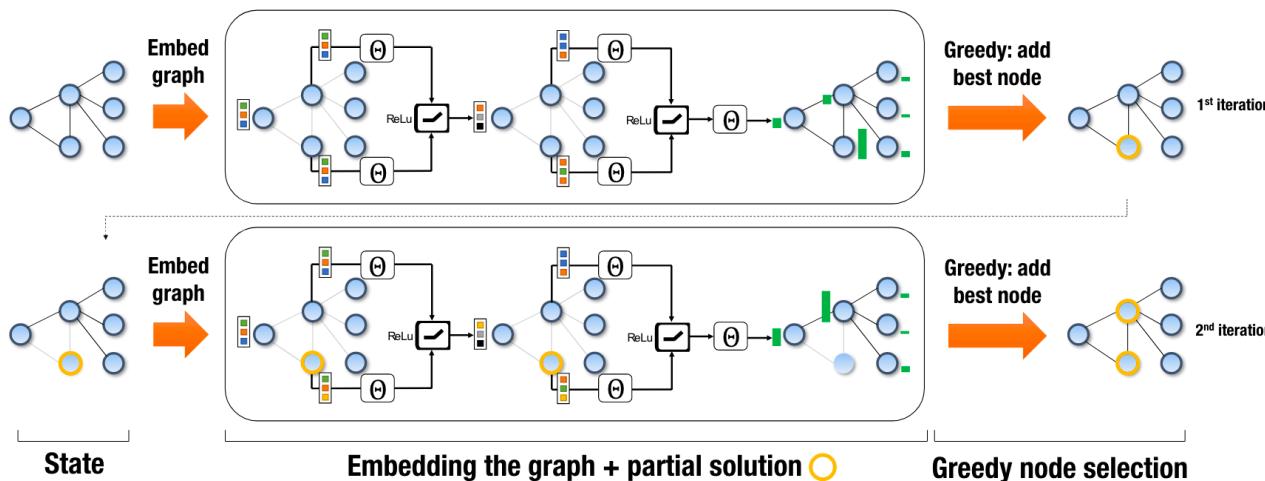
Cannot generalise to
more nodes



Example: Learn to Solve Minimum Vertex Cover

- **Problem:** Find the minimal subset of vertices to cover all the vertices in a graph
- **Design:**
 - Input: graph embedding (graph neural network)
 - Output: 0/1 for each vertex
 - Model: graph neural network
 - Reinforcement learning: learn a policy to add nodes

Can generalize
to different
graphs



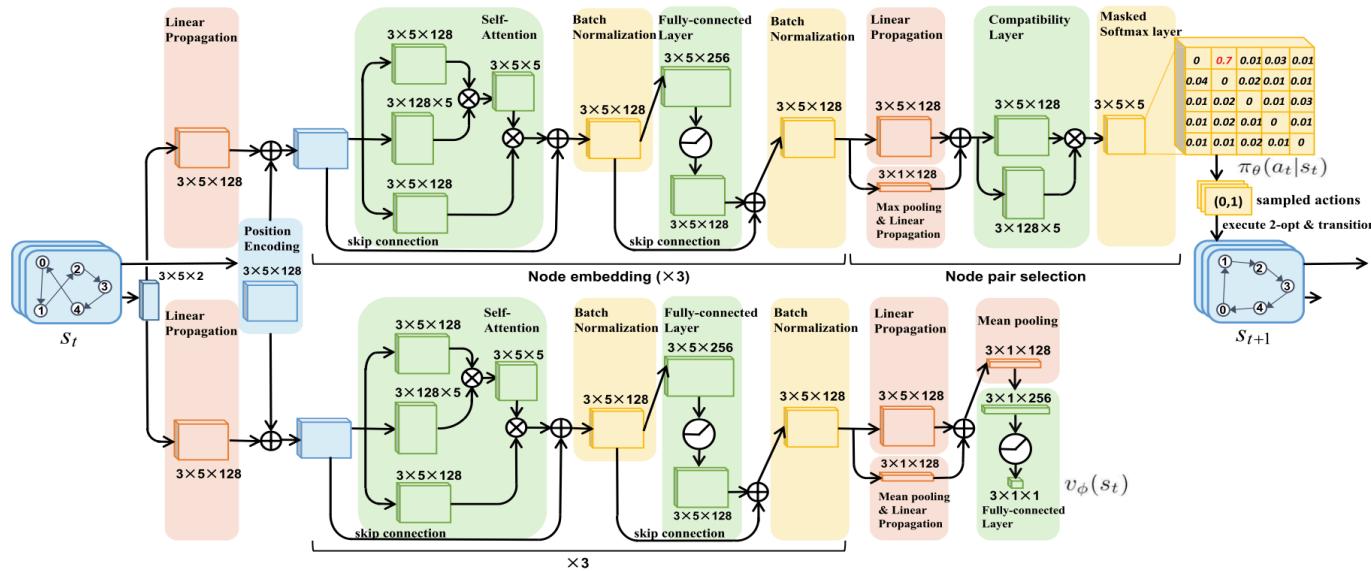
Learn to Search

- Machine learning to make design decisions of search algorithms
 - Select the next branch during branch and bound
 - Select the next neighbouring solution during local search
 - Select the next crossover/mutation operator in an EC algorithm



Example: Learn 2-opt for Routing

- Iterative improvement process
 - Solution -> 2-opt -> Solution -> 2-opt -> ...
 - Learn a **policy** to select a **pair of nodes** to conduct 2-opt
 - Reinforcement learning process

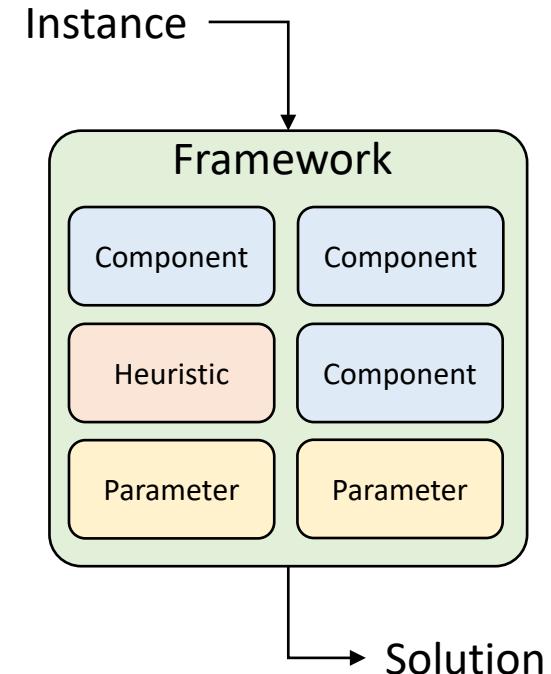




Evolutionary Computation to Learning Combinatorial Optimisation Heuristics

EC to Learn CO Heuristics

- Define a **space** of heuristics
 - A **framework** for the heuristic to work within
 - Input: a CO problem instance
 - Output: a solution
 - **Representation/Encoding** for heuristics
- Define the **evaluation** of heuristics
 - **Training data**
 - **Performance metric** (e.g., effectiveness, efficiency, complexity)
 - Black-box, non-differentiable, discrete
- Design **search algorithms by EC**
 - Good at black-box optimisation
 - Handle non-differential space
 - Multi-objective optimisation



EC to Learn CO Heuristics

Preprocessing

Design a **framework** Γ parameterized with the CO heuristic



Encoding

Decide the **space** \mathcal{H} of all possible CO heuristics in consideration



Data:

Determine a **training set** P (problem instances) to train the heuristics



Objective:

Design the training **objective function** $f(h, P, \Gamma)$

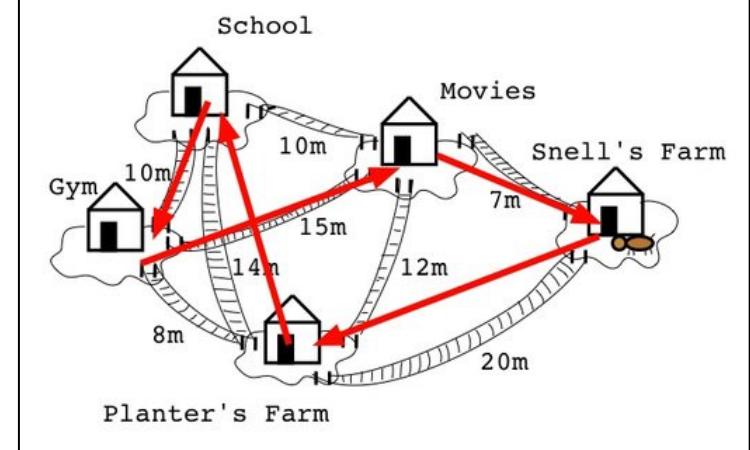
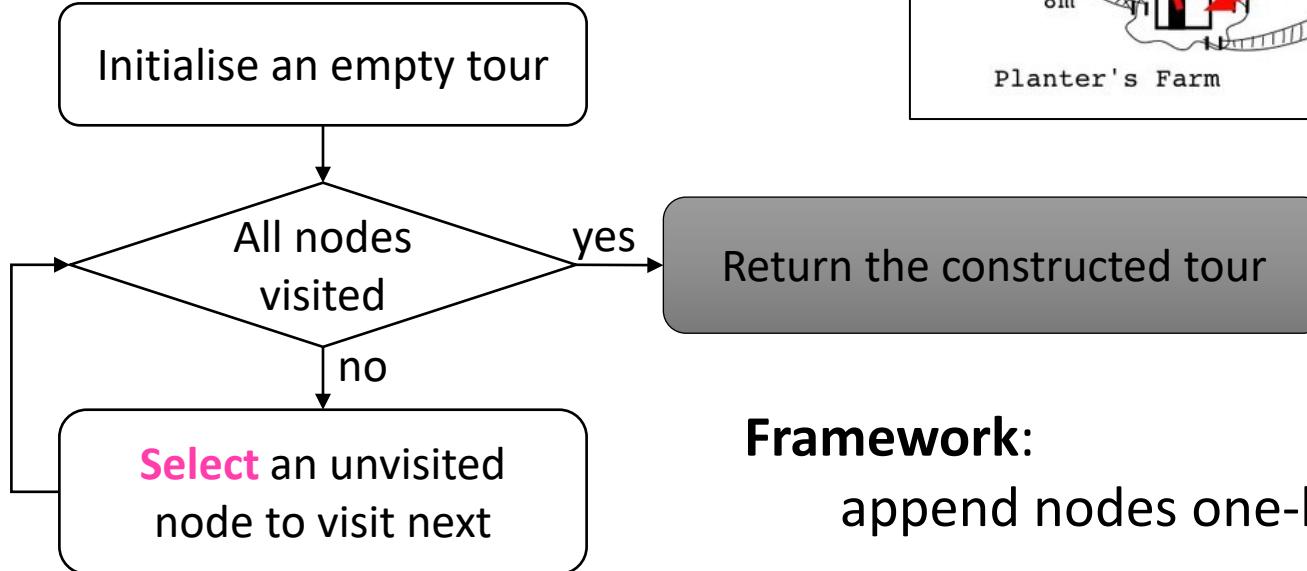


Optimisation:

Find the best heuristic $h^* = \arg \min_{h \in \mathcal{H}} f(h, P, \Gamma)$



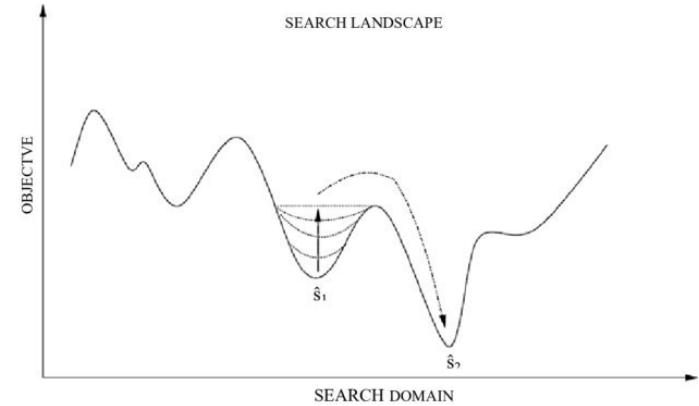
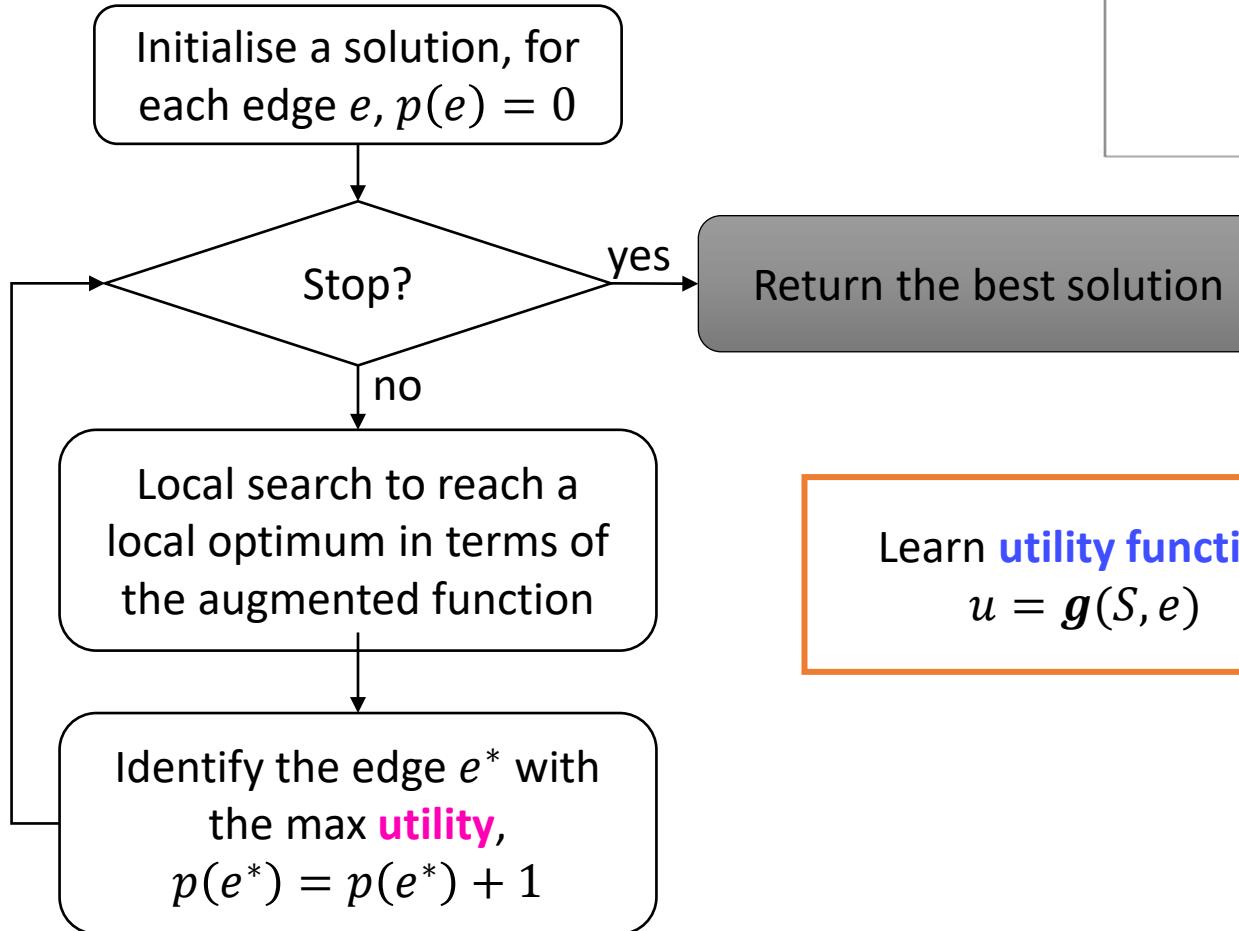
Example 1: TSP



Learn heuristic for **node selection**
 $x_{next} = h(G, X_{partial})$



Example 2: GLS for VRP



Learn **utility function**
 $u = g(S, e)$

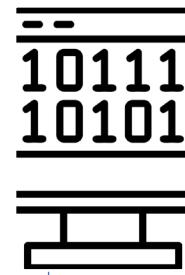


Genetic Programming for Heuristic Learning



Preprocessing

- Through **prior knowledge** and intuition
- A **simulation** to output a solution



Encoding

- **Trees**
- **Sequence** of instructions



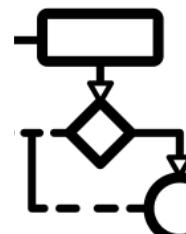
Objective function

- Normalised objective values of the obtained solutions



Data

- A wide range of **problem instances**
- Generalisation and efficiency (problem size, #instances) are important



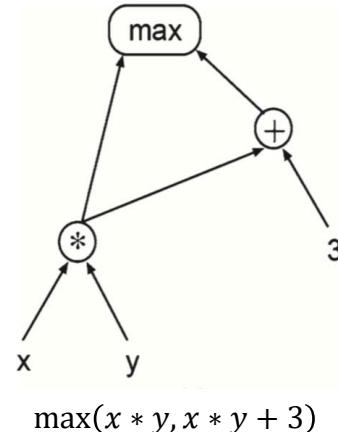
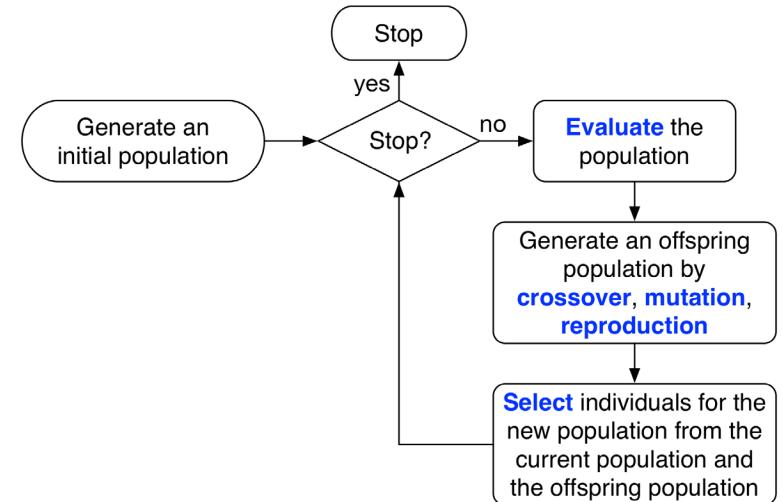
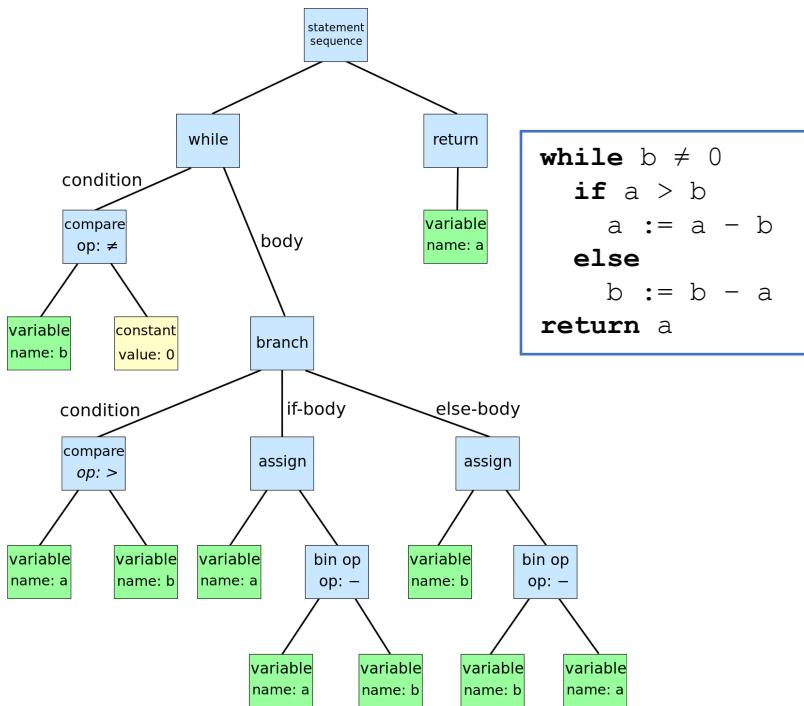
Optimisation

- Genetic programming



Genetic Programming (GP)

- A type of **evolutionary algorithm**
 - Evolve **computer programs** rather than solutions
- **Representation** of computer programs
 - Tree-like, graph-like, linear, ...



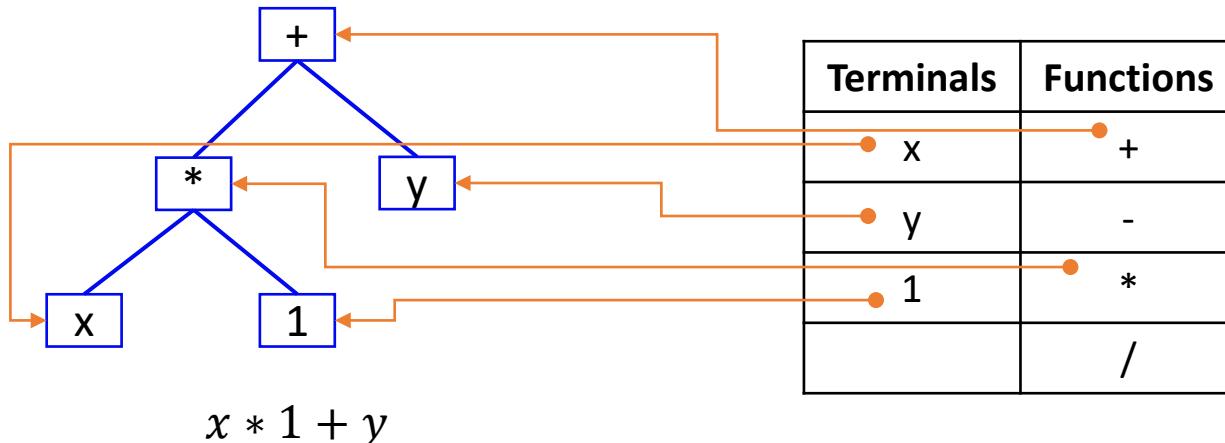
```

r[3] = r[1] / 1.3;
r[1] = r[2] * -5.5;
r[0] = if r[1] < 0 then
       r[0] else r[3];
  
```



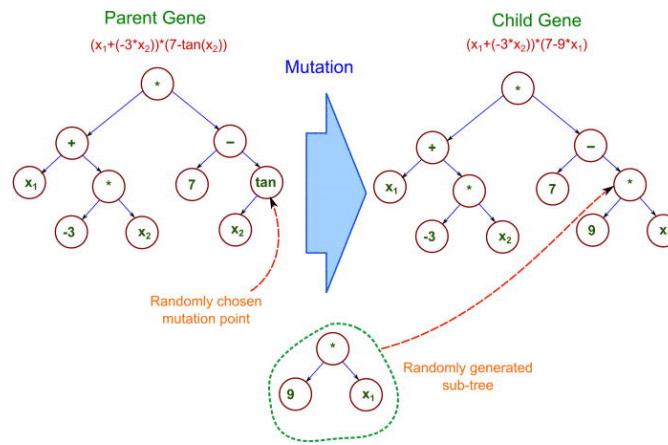
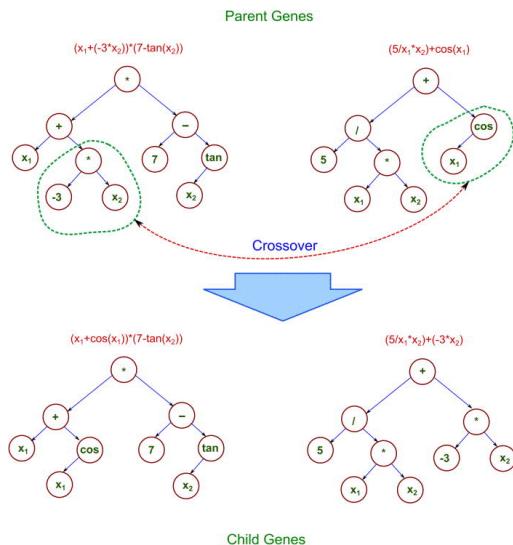
Genetic Programming (GP)

- Individual generation (Tree-based representation)
 - Terminal set: inputs of the program and constants, no argument, form the leaf nodes,
 - Function set: operators to the inputs and intermediate results of the program (e.g., +, -, max, ...), form the non-leaf nodes
- Start from the root node
- For each node, randomly sample from the terminal/function set
 - If sampling from the terminal set, then stop this branch
 - If sampling from the function set, create the child nodes, and recursively sample the child nodes

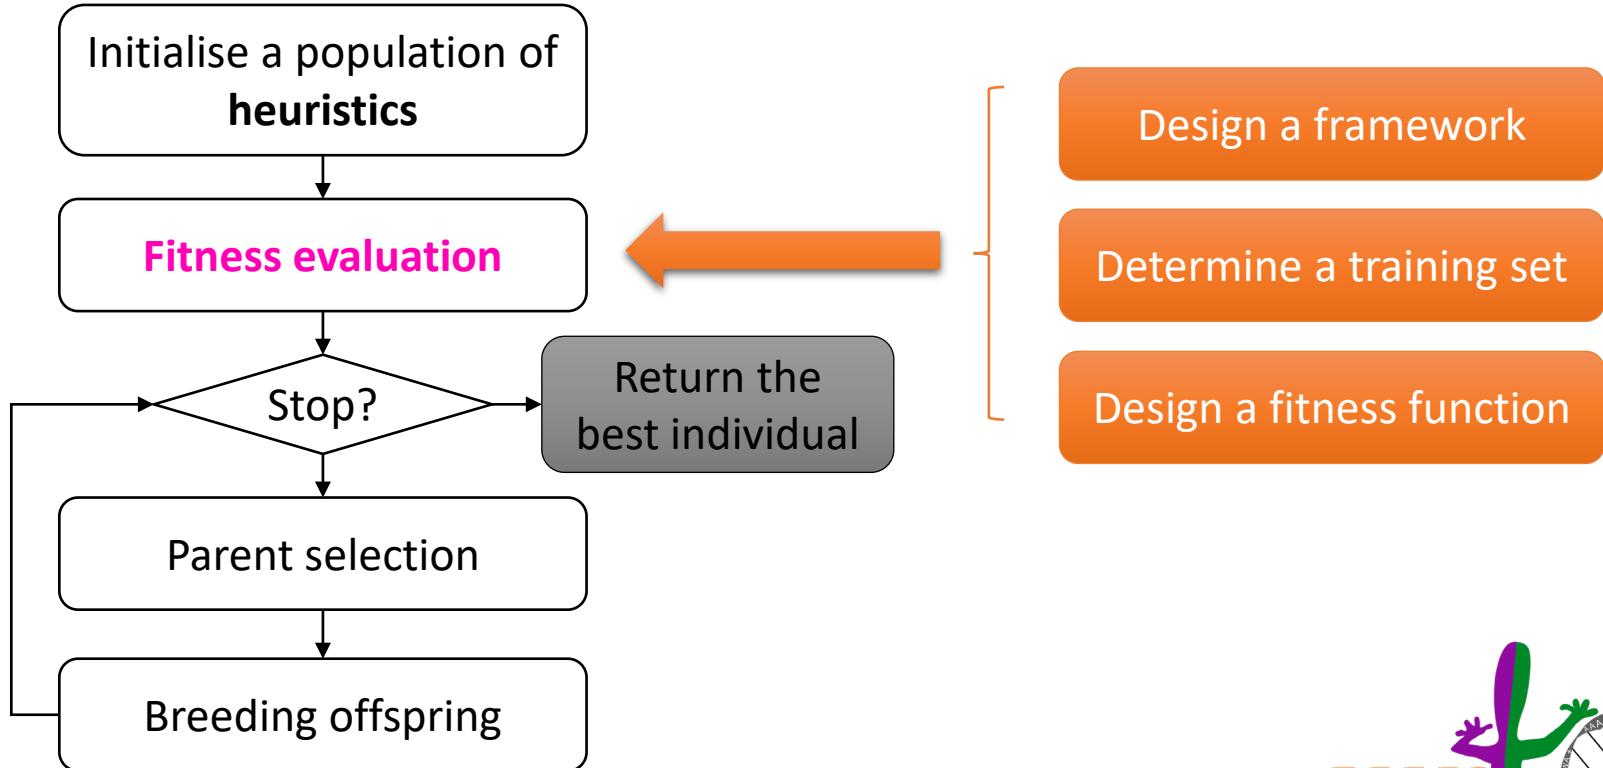


Genetic Programming (GP)

- **Genetic Operators** (tree-based)
 - **Crossover**: randomly select a sub-tree from each parent, and swap them
 - **Mutation**: randomly select a sub-tree from the parent, and replace the sub-tree with a randomly generated sub-tree (e.g., grow)
 - **Reproduction**: copy the parent directly



GP to Learn CO Heuristic



The **framework** with the **heuristic** obtains/returns a **solution** to a given **instance**



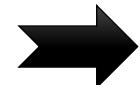
Fitness Evaluation

Heuristic



Framework

Training	Ref value	Heuristic	Framework
Instance 1	120	Heuristic	Framework
Instance 2	240	Heuristic	Framework
Instance 3	180	Heuristic	Framework

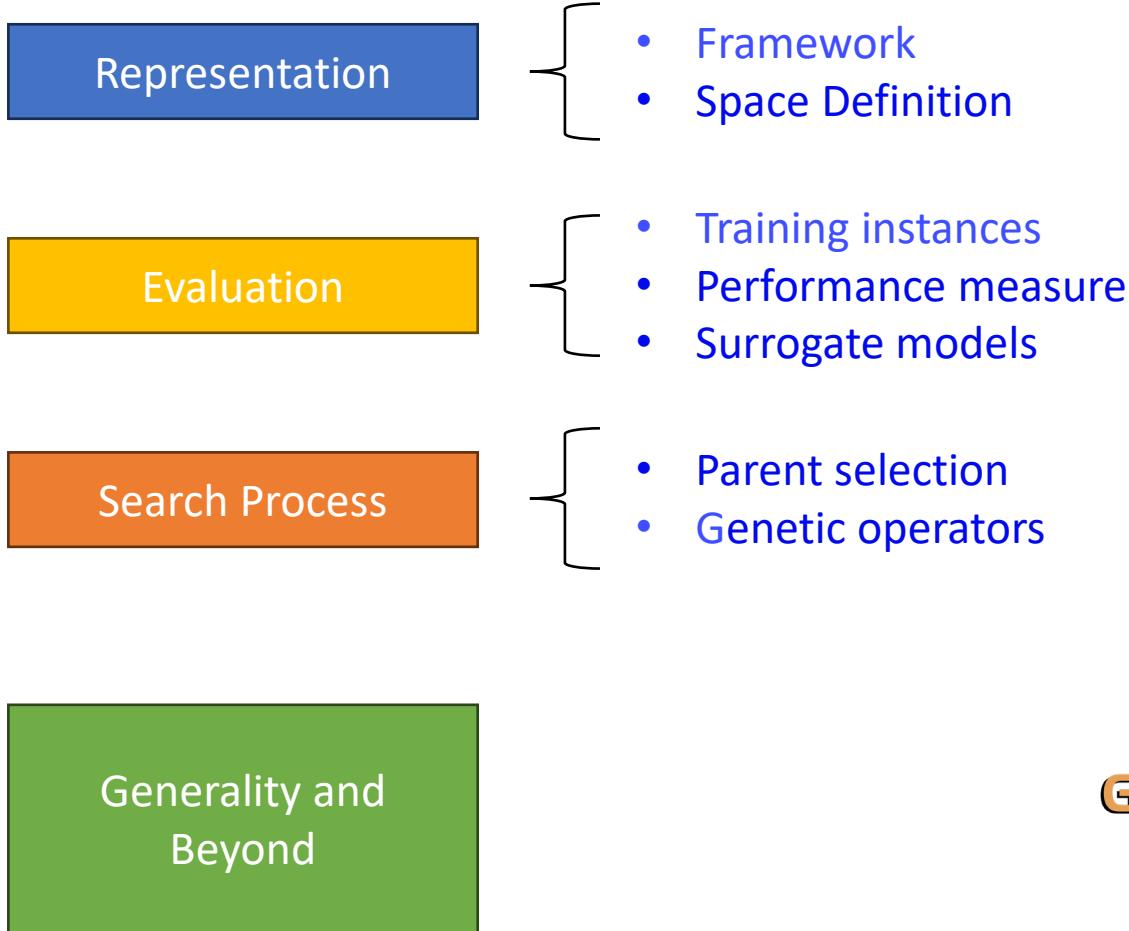


Solution	Obj
S1	140
S2	250
S3	200

$$fitness = \frac{\left(\frac{140}{120} + \frac{250}{240} + \frac{200}{180} \right)}{3} = 1.11$$



Design Questions



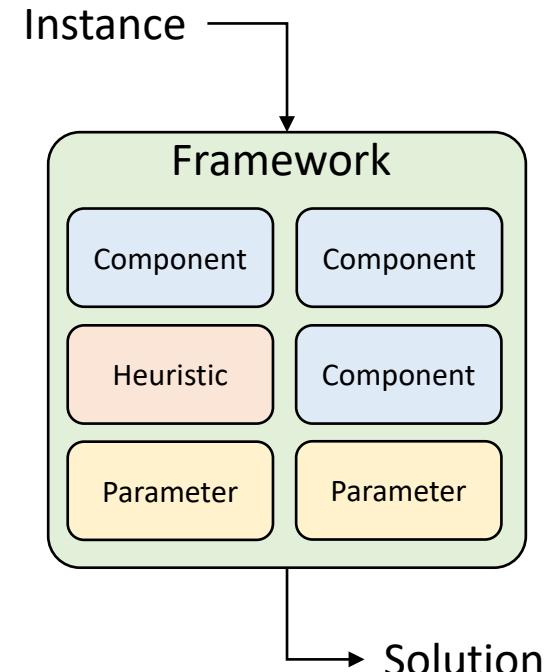
Representation

- **Framework**

- **Input:** Instance
- **Output:** Solution
- **Parameters:** Heuristic to be learned
 - Constructive heuristic
 - Improvement heuristic / operator
 - Parameters

- **Space Definition** (Set of all possible heuristics)

- Features / Terminals
- Functions / Non-terminals
- Model structure (tree-based, linear-based, ...)
- Possible combinations (Grammar)



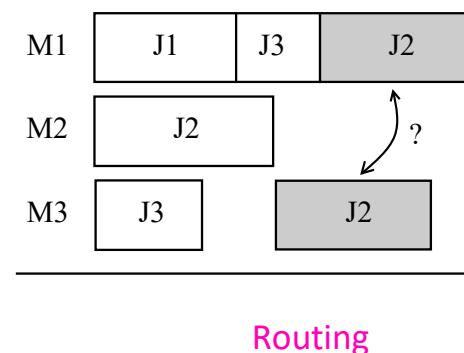
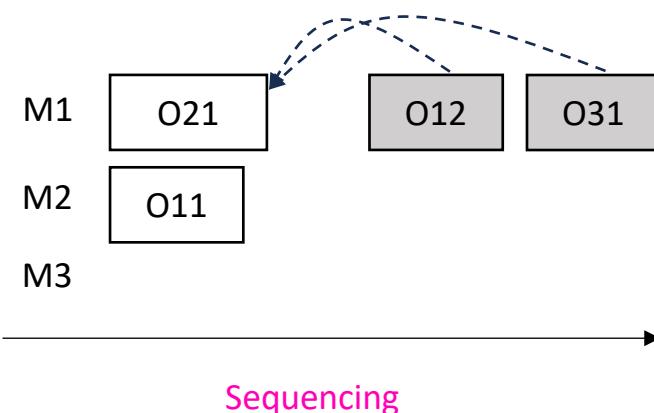
Framework

- The framework is usually designed based on **intuition** and **domain knowledge**
 1. Take an existing algorithm: everything is **manually designed** and **hard coded**
 2. Identify the **component(s)** of the existing algorithm to be learned
 3. **Parametrise** the identified component(s)
- **Example**
 - Learn flexible job shop scheduling heuristics



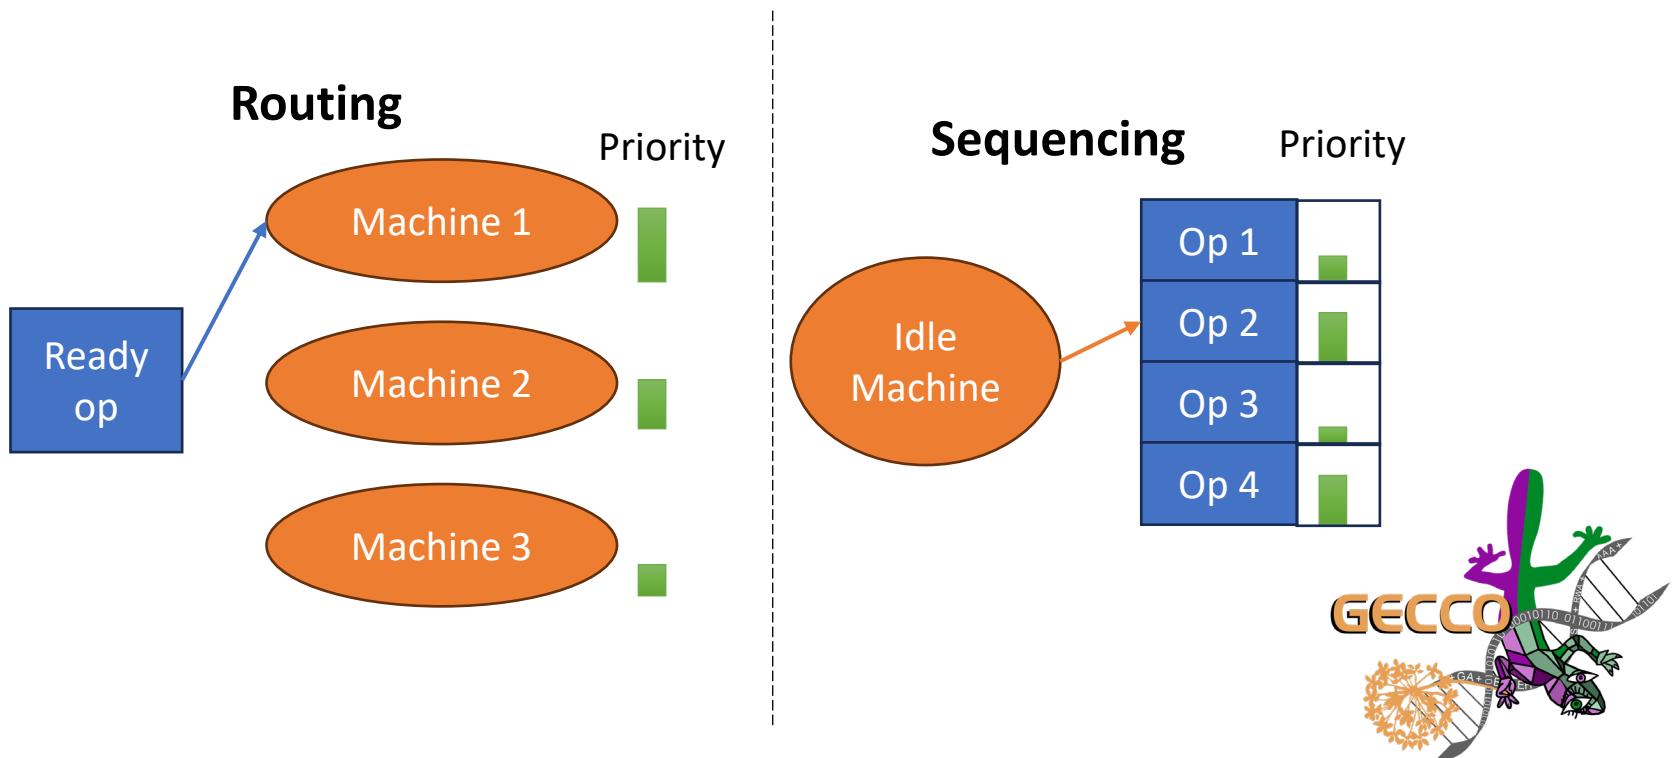
Flexible Job Shop Scheduling Heuristics

- Process set of **jobs** by a set of **machines**
- Each job has a sequence of **operations**
- Each operation can be processed by multiple candidate **machines**
- **Routing** decision: select a machine for an operation
- **Sequencing** decision: select the next operation from a machine's queue



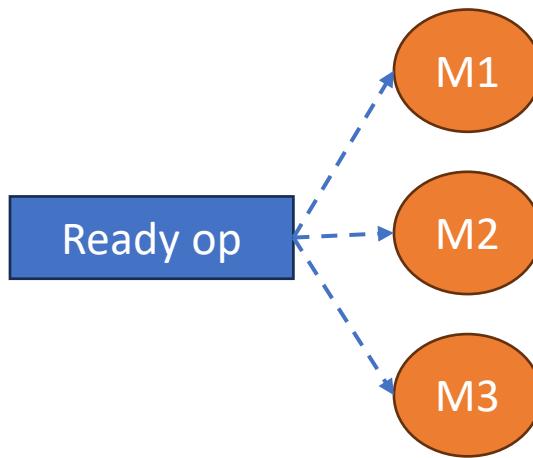
Framework 1: non-delay

- Select machine as soon as an operation is ready
- Select operation as soon as a machine becomes idle

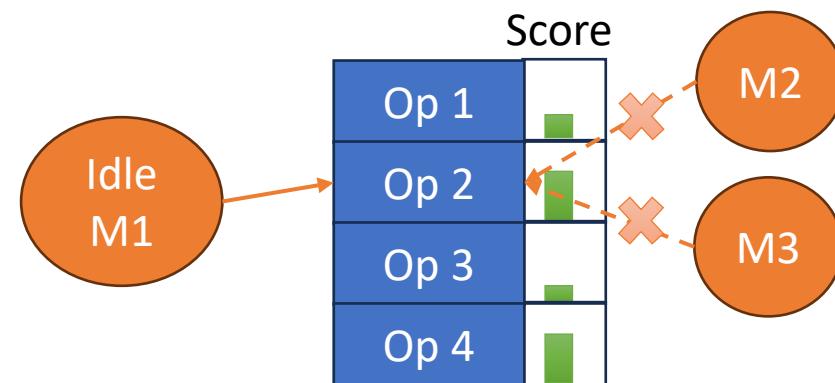


Framework 2: delayed-routing

- Routing (machine selection) is delayed until at least one candidate machine becomes idle



No Routing:
no machine is idle



Sequencing
(+ delayed routing)



Space Definition: Terminals

- Consider all the features that might affect the decisions
- But this will lead to a **huge GP search space**
- Can alleviate by **feature/terminal selection**

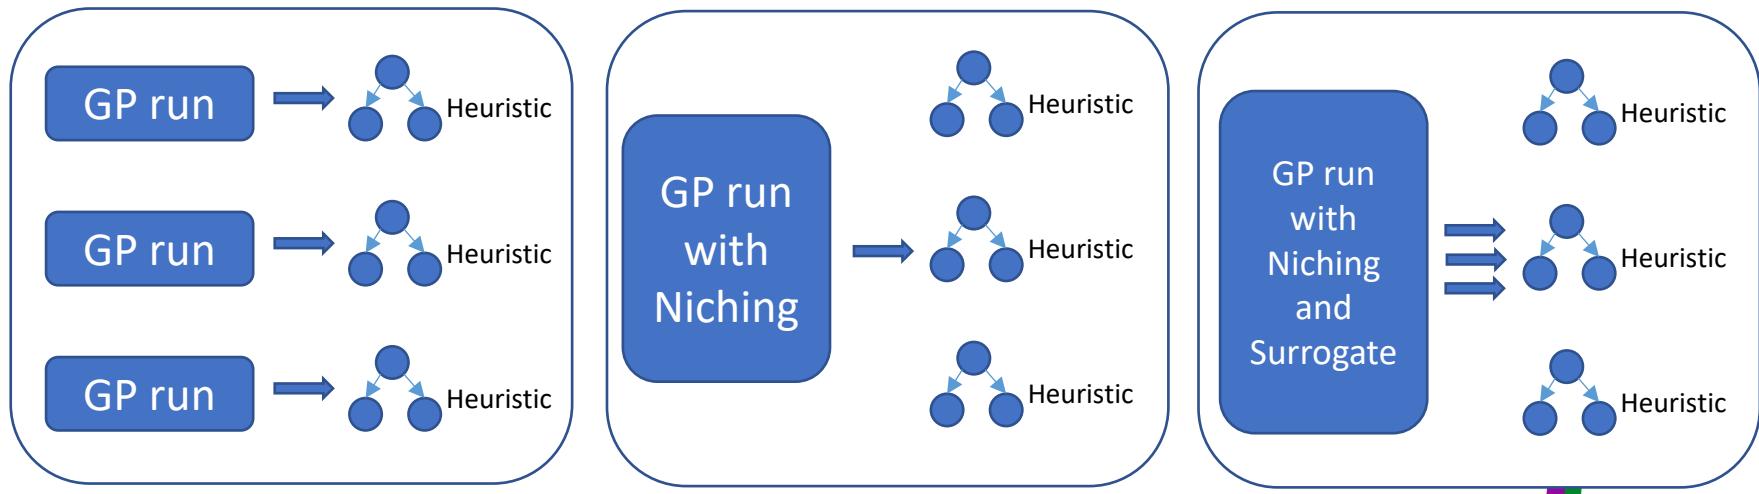
- **Q:** How do we know which features are important?
- **A:** If a feature is used by many different good heuristics

- But we don't have many different good heuristics



Space Definition: Terminals

1. Collect a diverse set of good heuristics
2. Calculate the importance of each terminal to each heuristic
3. Select terminals by aggregated importance



[Intuitive]: Several independent GP runs, gather the best-of-run heuristics

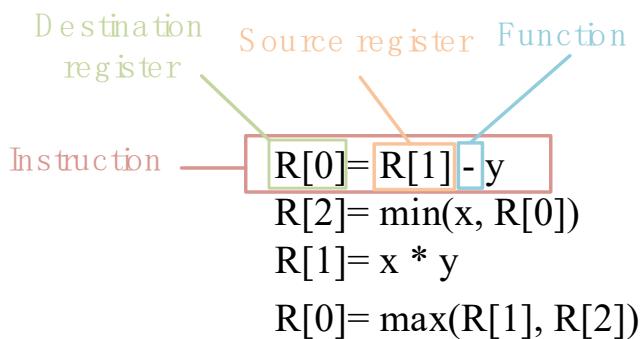
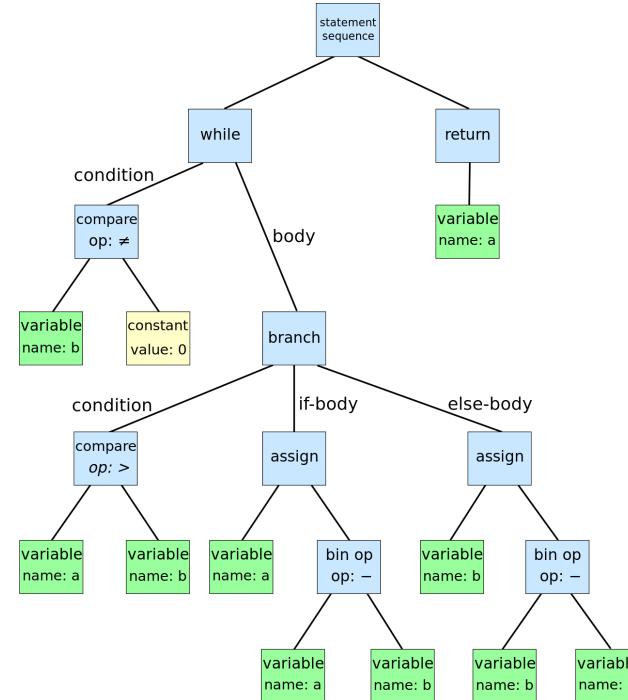
[Better]: A single GP run with niching, gather the diverse local optimal heuristics

[Best]: A single GP run with niching, and surrogate evaluation

- Yi Mei, Mengjie Zhang, and Su Nyugen. Feature selection in evolving job shop dispatching rules with genetic programming. In Proceedings of the Genetic and Evolutionary Computation Conference (GECCO), pages 365–372. ACM, 2016.
- Yi Mei, Su Nguyen, Bing Xue, and Mengjie Zhang (2017). An efficient feature selection algorithm for evolving job shop scheduling rules with genetic programming. IEEE Transactions on Emerging Topics in Computational Intelligence, 1(5), 339-353.

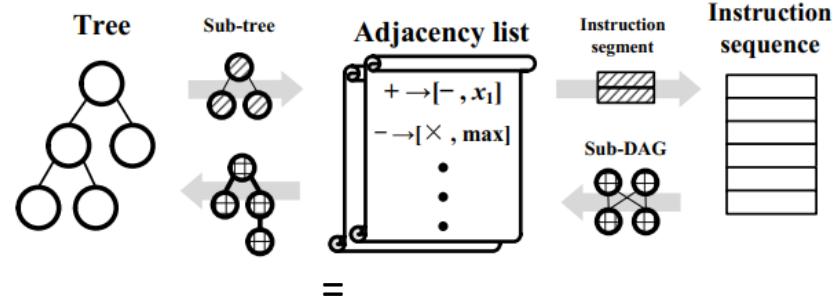
Space Definition: Model Structure

- Tree-based structure
 - Abstract syntax tree of program
- Linear-based structure
 - A sequence of register-based instructions
 - Can represent a Directed Acyclic Graph (DAG)
 - Easy reuse of building blocks

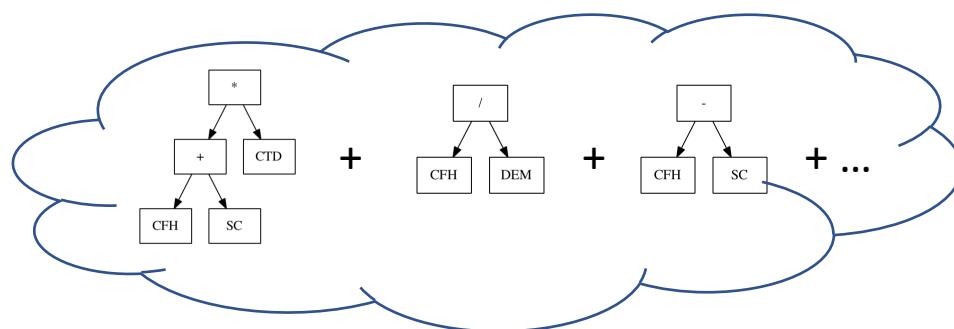


Space Definition: Model Structure

- Using tree-based and linear-based representation **simultaneously**
 - Multiple sub-populations, each with a different representations
 - Interact with crossover
 - Adjacency list as intermediate

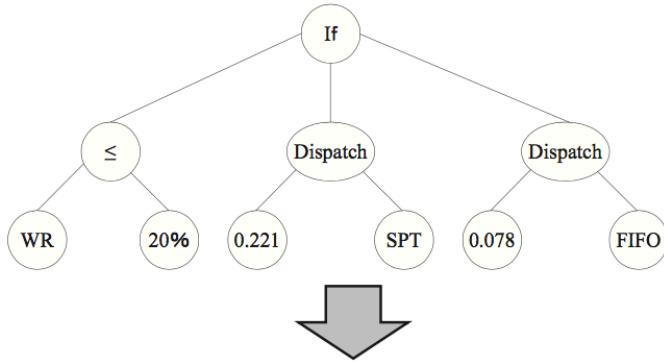


- Multi-tree / Ensemble representation
 - Each tree stands for a different heuristic (e.g., routing and sequencing rules)
 - Trees collaborate to make decision (e.g., voting)



Space Definition: Grammar

- Restrict how features/terminals can be combined
 - Smaller search space, better interpretability
 - Disconnected space, may lose effectiveness
- Example: grammar for meaningful dispatching rules
 - Recursive depth-first sampling, generate a tree



```
If (workload ratio is less than or equal to 20%)  
    Use the SPT rule with non-delay factor of 0.221  
Else  
    Use the FIFO rule with non-delay factor of 0.078
```

```
Start ::= <action>  
<action> ::= <if> | <dispatch>  
<if> ::= if <attributetype><op> <threshold>  
then <action> else <action>  
<op> ::= ≤ | >  
<attributetype> ::= WR | MP | DJ | CMI | CWR | BWR  
<threshold> ::= 10% | 20% | 30% | 40% | 50% | 60% | 70% | 80% | 90% | 100%  
<dispatch> ::= assign <nondelayfactor> assign <rule>  
<nondelayfactor> ::= uniform[0,1]  
<rule> ::= FIFO | SPT | LPT | LSO | LRM | MWKR | SWKR |  
MOPR | EDD | MS | WSPT
```



Space Definition: Grammar

- Linear GP can also have grammar!
 - **Module** Context-Free Grammar for sequence of instructions

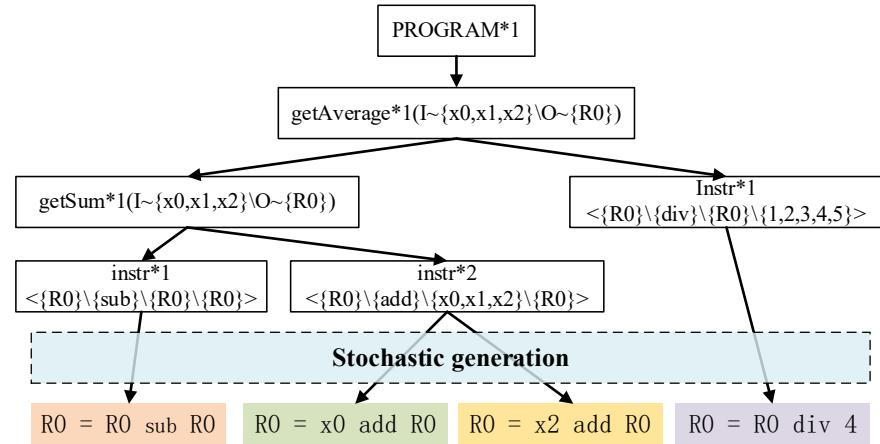
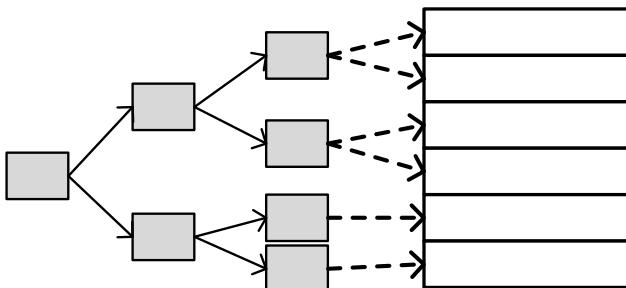
Grammar rules based on domain knowledge

```
defset FUNS {add,sub,mul,div};
defset REG {R0,R1,R2,R3,R4};
defset INPUT {x0, x1, x2};
defset CONSTANTS {1,2,3,4,5};

getSum(I\O) ::= <0\{sub\}\0>::<0\{add\}\I\0>*5;
getAverage(I\O) ::= getSum(I~I\0~0)::<0\{div\}\0\{CONSTANTS>;
getVariance(I\O) ::=getAverage(I~I\0~{R4})::(<{R1,R2,R3}\{sub\}\{INPUT\}\{R4>;
::<{R1,R2,R3}\{mul\}\{R1,R2,R3\}\{R1,R2,R3>)*3::getAverage(I~{R1,R2,R3}\0~0);
PROGRAM ::= getSum(I~{INPUT}\0~{R0}) | getAverage(I~{INPUT}\0~{R0})
| getVariance(I~{INPUT}\0~{R0});
```

Derivation tree
based on grammar

LGP program based
on derivation tree



Evaluation

- **Training instances**
 - A wide range of training instances is needed for generalization
 - However, too many training instances will slow down the evaluation and waste resources
 - Generalisation vs Efficiency
- **Performance metrics**
 - Effectiveness (e.g., quality of obtained solutions)
 - Efficiency (computational complexity)
 - Interpretability
 - ...
- **Surrogate models**



Training Instances

- A wide range of training instances + instance rotation
 - Similar to **batch learning**
 - **Split** the training instances into small subsets (batches)
 - At each generation, use a **different subset** (same for all individuals in the population)

Training set



Instance rotation



Gen 1

Gen 2

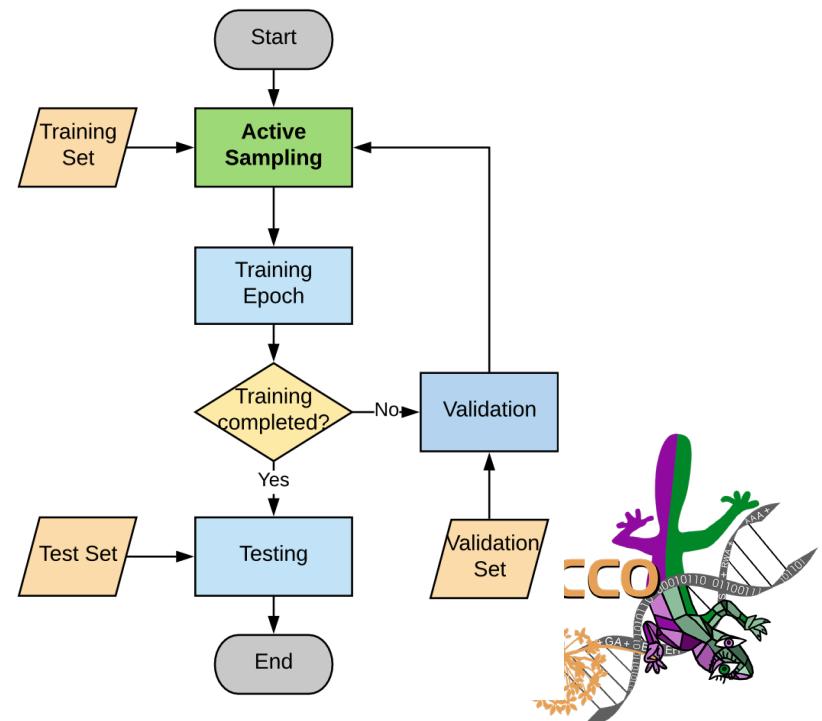
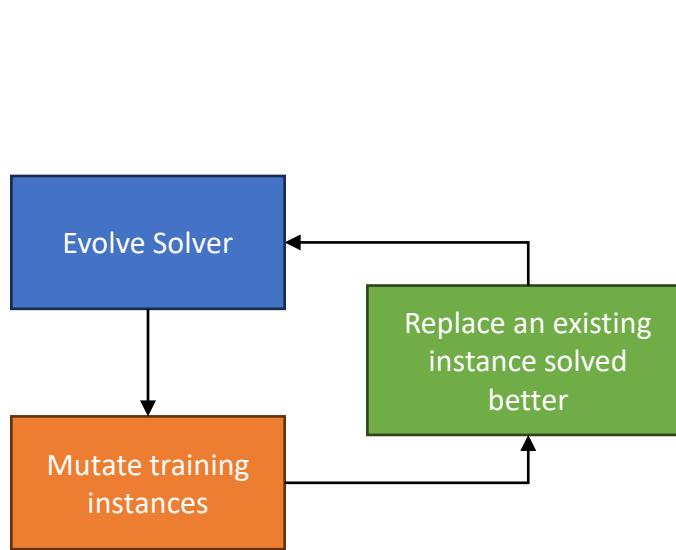
Gen 3

.....



Training Instances

- Intelligently select/sample the best training instances
 - Remove the instances that have been solved well
 - Add/Retain the instances that are not solved
 - Focus on the instances that best distinguish good/bad heuristics



- Tang, K., Liu, S., Yang, P., & Yao, X. (2021). Few-shots parallel algorithm portfolio construction via co-evolution. *IEEE Transactions on Evolutionary Computation*, 25(3), 595-607.
- Karunakaran, D., Mei, Y., Chen, G., & Zhang, M. (2019, June). Active sampling for dynamic job shop scheduling using genetic programming. In *2019 IEEE Congress on Evolutionary Computation (CEC)* (pp. 434-441). IEEE.

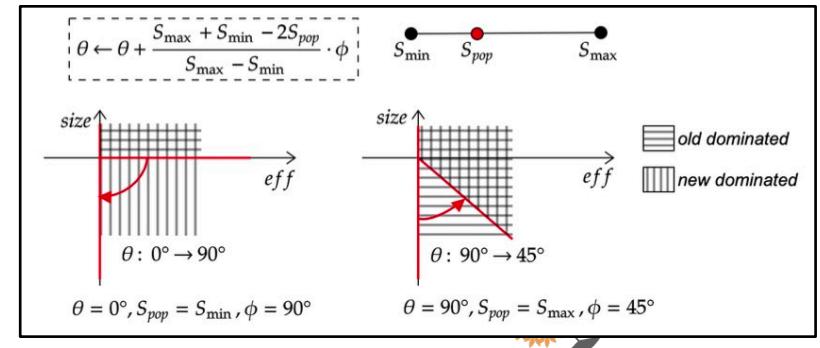
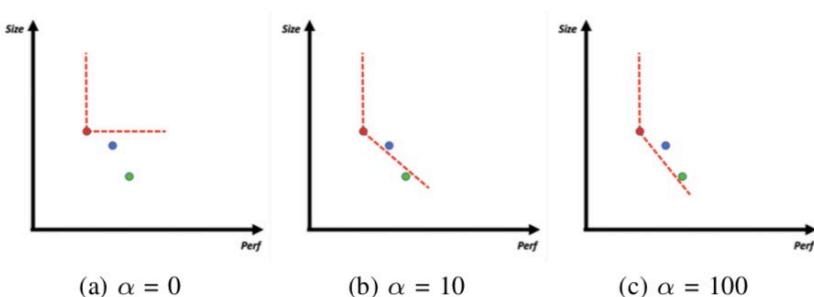
Performance Metrics

- Average/Normalised solution quality obtained over the training/validation set (Effectiveness)
- For iterative improvement frameworks, **when to stop?**
 - After converge, balance between effectiveness and efficiency
 - During training, as early as possible if can distinguish good and bad heuristics
- For stochastic framework/heuristic, need multiple runs for statistical significance (efficiency issue)
 - Learn deterministic algorithms (e.g., ϵ greedy)
- Model size/complexity, interpretability
- **Multi-objective** optimisation



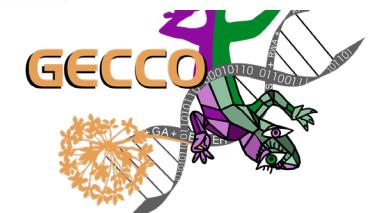
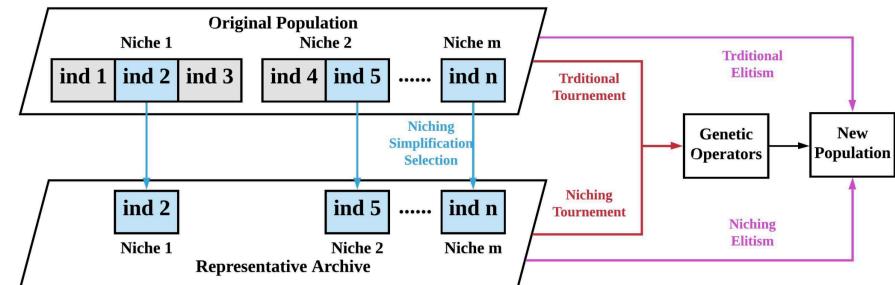
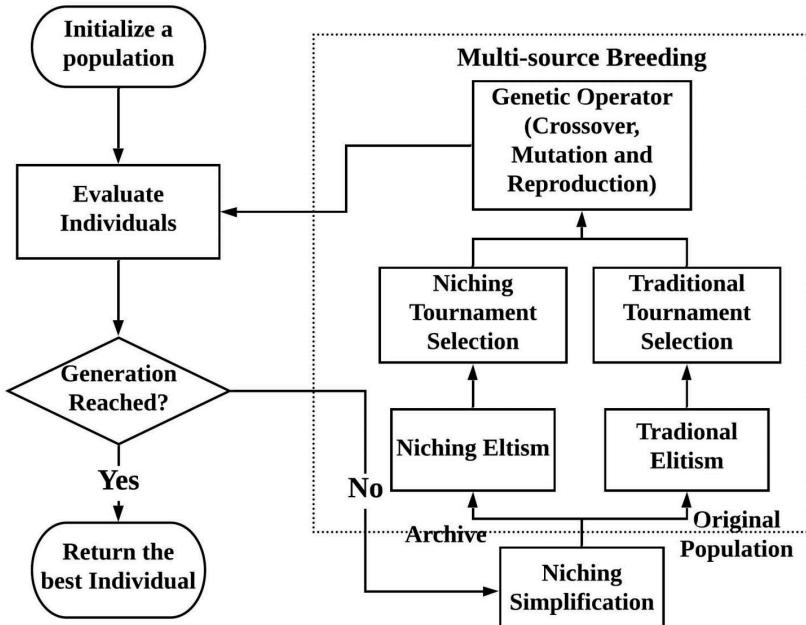
Effectiveness vs Interpretability

- MOGP: $f_1 = \text{effectiveness}$; $f_2 = \text{model size} (\text{interpretability proxy})$
- It is hard to balance effectiveness (e.g., accuracy) and size during the MOGP search
- If not evolve properly, the population can be easily biased to small but bad individuals, and lose exploration ability
- Use α -dominance to adjust the balance between effectiveness and size
 - $\alpha = 0$: normal dominance relationship
 - $\alpha = \infty$: single objective with only effectiveness
- Adapt α during the search



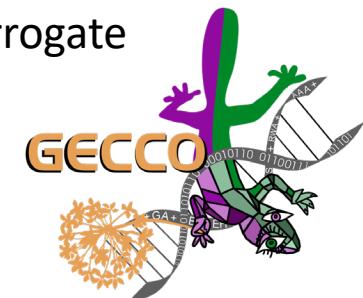
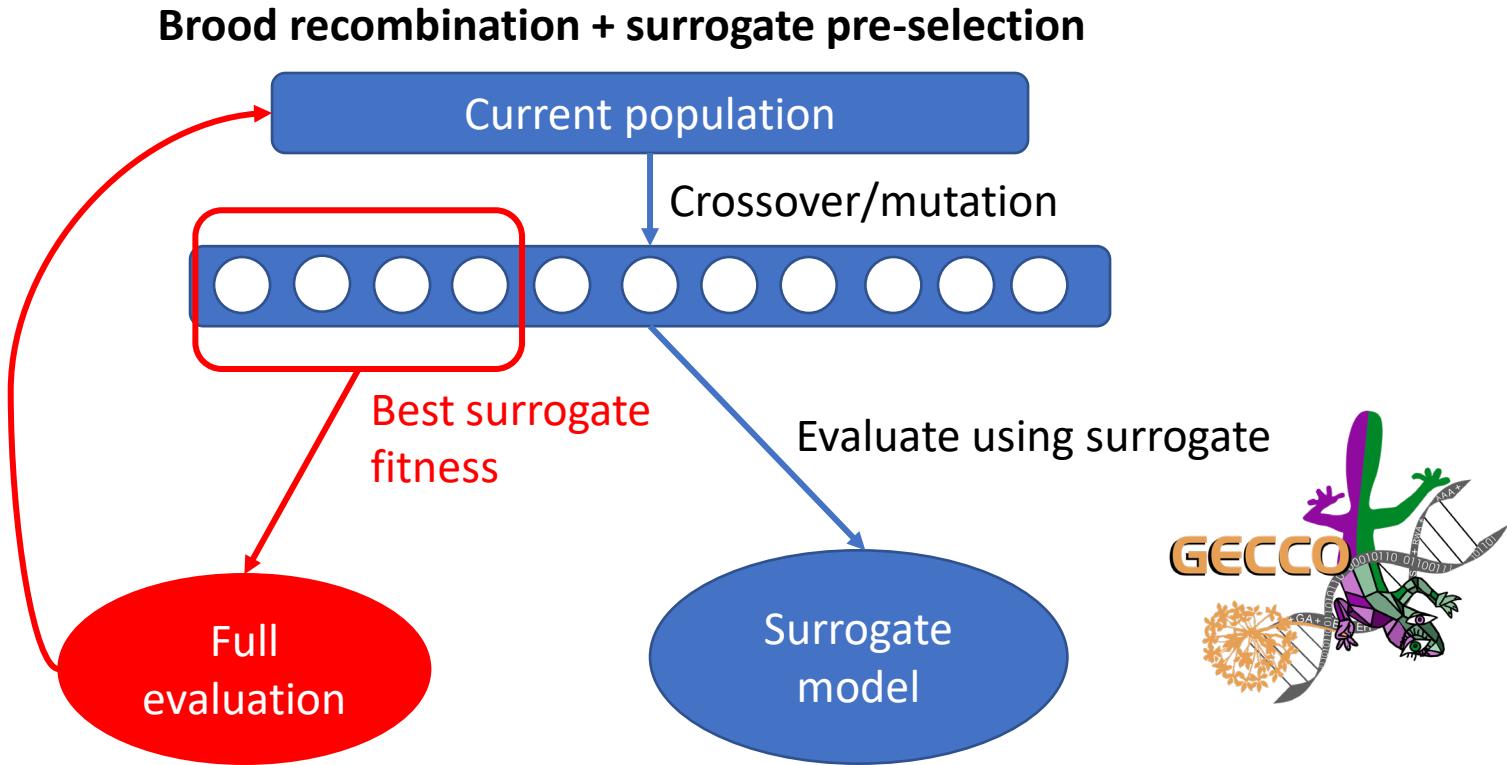
Effectiveness vs Interpretability

- Niching GP
- Group individuals into niches, each niche contains individuals with the same fitness value
- Niching selection: select the smallest individual from each niche
- Multi-source breeding



Surrogate Models

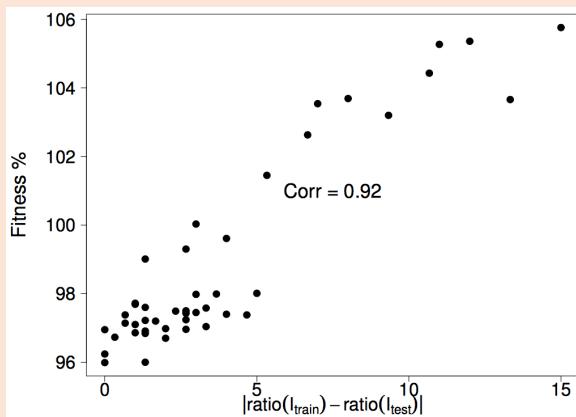
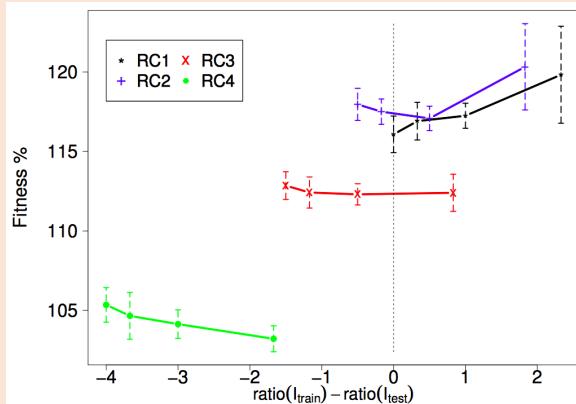
- Brood recombination: generate a large number of offspring
- Surrogate pre-selection: use surrogate to predict fitness and select the best parts



Surrogate Models

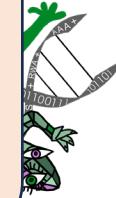
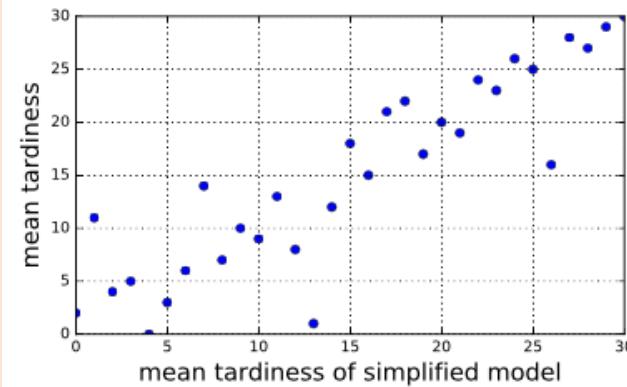
- Small instances with careful **configurations**, faster and correlated

Static JSS



Dynamic JSS

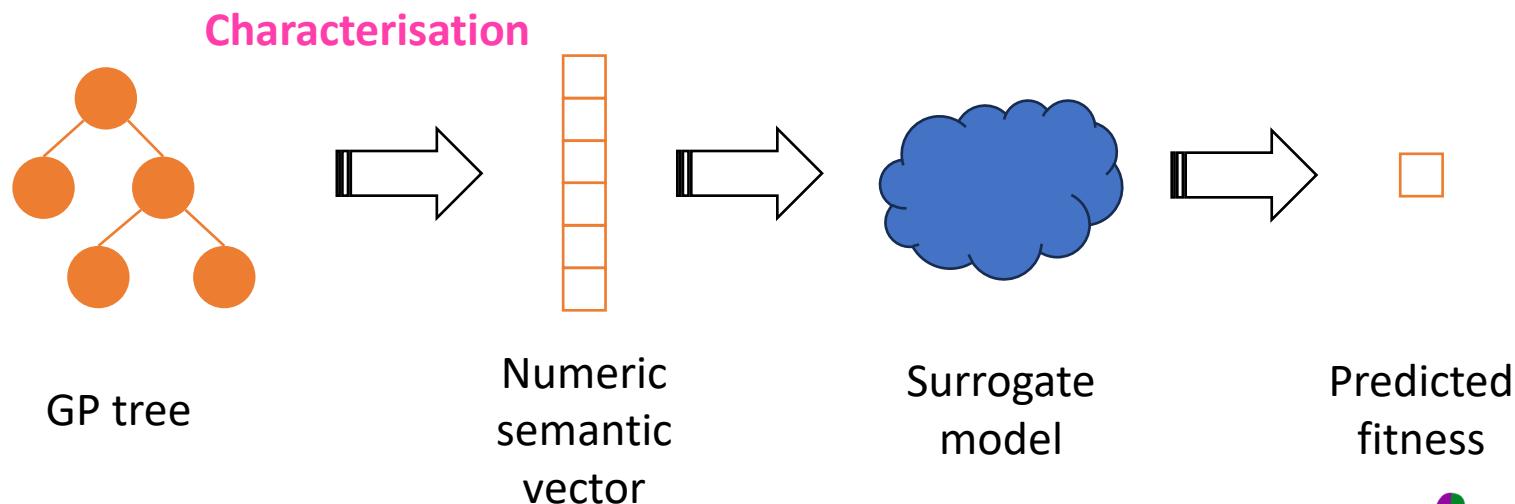
Configuration	Original	Surrogate
No. Machines	10	5
No. Jobs	5000	500
No. Warmup Jobs	500	100
Min Ops	2	2
Max Ops	14	7



- Yi Mei, Mengjie Zhang, "A Comprehensive Analysis on Reusability of GP-Evolved Job Shop Dispatching Rules," *IEEE World Congress in Computational Intelligence (WCCI)*, Vancouver, Canada, 2016.
- Nguyen, S., Zhang, M. and Tan, K.C., Surrogate-Assisted Genetic Programming With Simplified Models for Automated Design of Dispatching Rules. *IEEE Transactions on Cybernetics*, DOI 10.1109/TCYB.2016.2562674.

Surrogate Models

- Semantic/Phenotypic Characterisation
 - Convert a tree structure to a numeric vector



Surrogate Models

- An example in learning dispatching rules for dynamic job shop scheduling
 - A set of **decision situations** and a **reference rule**
 - For each decision situation, measure the difference between the reference rule and the characterised rule
 - Characterised by a **decision vector**

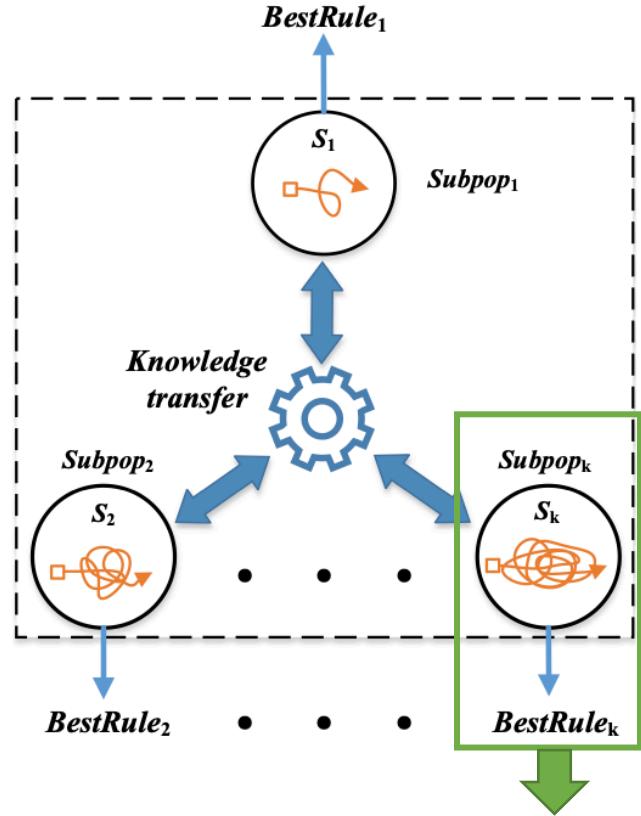
decision situation	attribute set s			ranking by reference rule	ranking by other rule	decision vector d
	s_1	s_2	s_3			
1	3	4	8	1	2	2
	7	6	15	2	1	
2	23	17	1	2	2	3
	8	9	3	3	1	
2	6	4	6	1	3	3
	:	:		:	:	
k	7	3	9	2	2	1
	4	8	6	1	1	

2
3
...
1



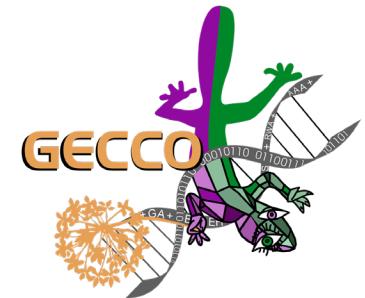
Surrogate Models

- One can have **multi-fidelity** surrogate models
 - Different problem size, model complexity, ...
 - Accuracy vs Speed
- Use **multiple sub-populations**
 - Each uses a different surrogate (**different fidelities**)
 - Sub-populations **transfer knowledge** by sharing parents to generate offspring
- One sub-population uses the original model (highest fidelity) and returns the final individual



Search Process

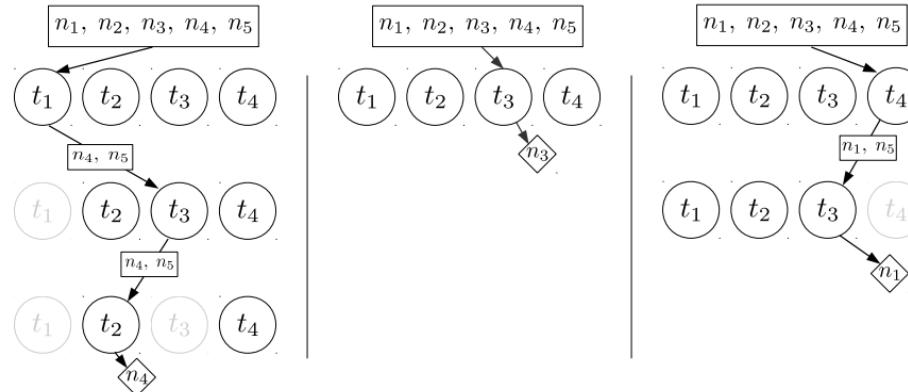
- Parent Selection
 - Commonly used: tournament selection
 - Lexicase selection for better diversity
- Genetic Operators



Lexicase Selection

- Find **specialists** instead of generalists
 - Multiple cases that compose the fitness function
 - Randomly iterate through cases
 - For each case, only keep the individuals with the **best case-fitness**
 - Stop if only one individual retains
- Case definition
 - Instances:

$$\text{fitness} = \sum_{i=1}^N (y_i - g(\mathbf{x}_i))^2 \quad \text{Case-fitness}$$
$$\text{fitness} = \sum_{i=1}^N \text{obj}(\Gamma(h, i)) \quad \text{Case-fitness}$$



- Helmuth, T., Spector, L., & Matheson, J. (2014). Solving uncompromising problems with lexicase selection. *IEEE Transactions on Evolutionary Computation*, 19(5), 630-643.
- Xu, M., Mei, Y., Zhang, F., & Zhang, M. (2023). Genetic programming with lexicase selection for large-scale dynamic flexible job shop scheduling. *IEEE Transactions on Evolutionary Computation*.

Genetic Operators

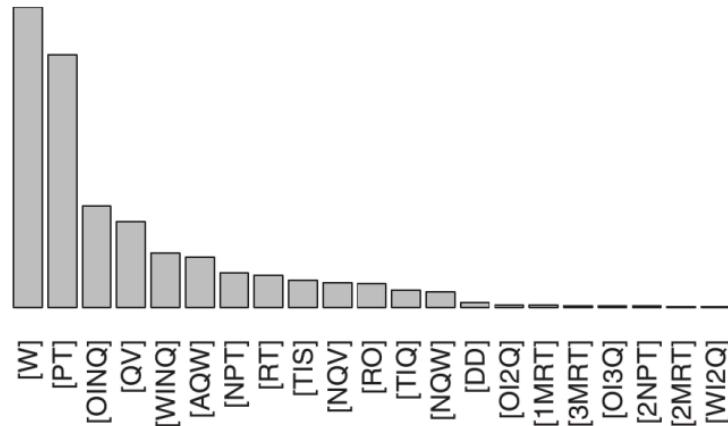
- Completely random
- **Adaptive** during the search process
 - Adjust the probability distribution of sampling nodes in mutation
 - In crossover, replace less important sub-trees with more-important sub-trees



Adaptive Mutation

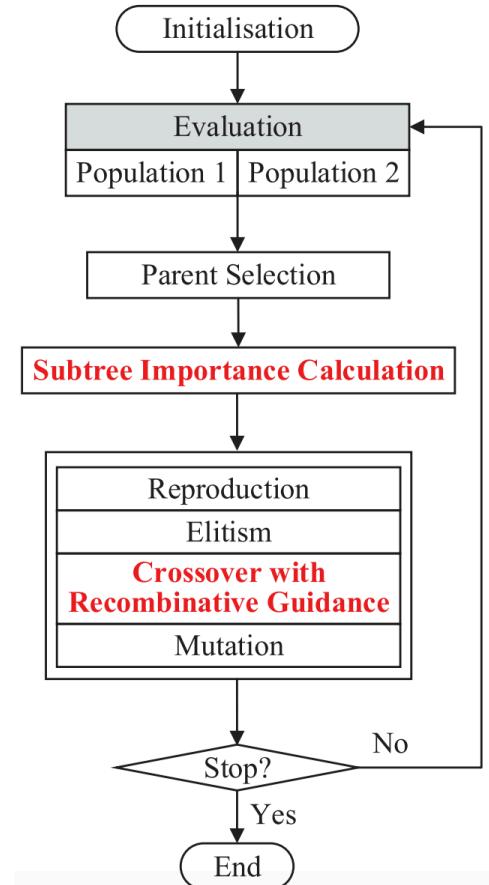
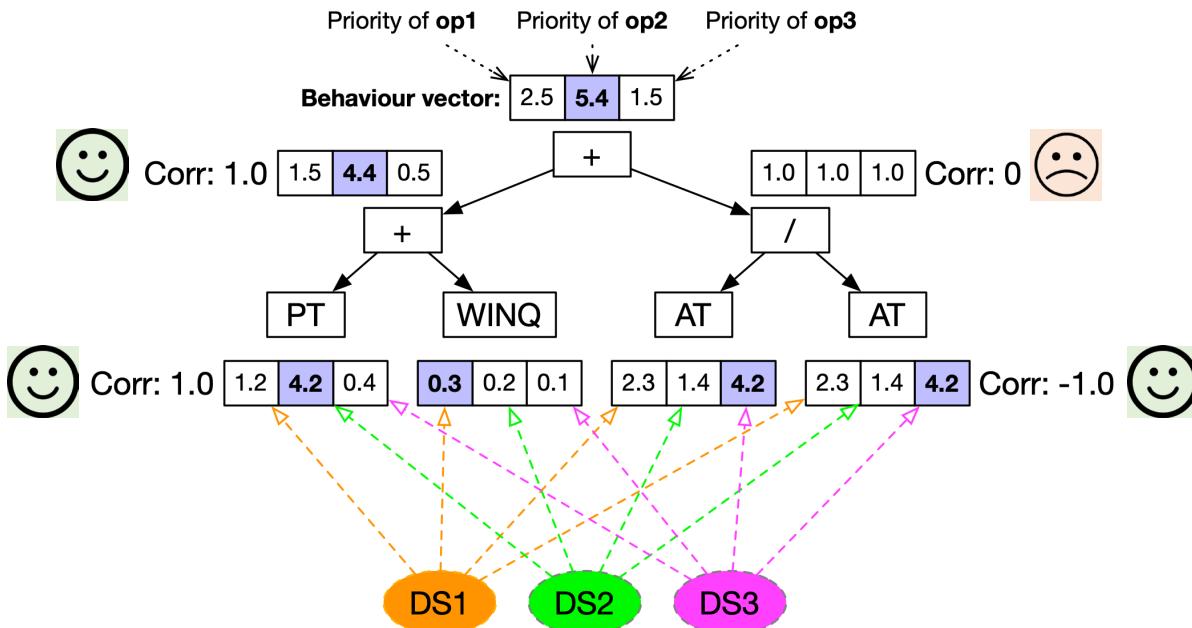
- Adjust probability of terminals based on historical occurrences
 - A terminal is sampled with a higher probability if it occurs more often in individuals with better fitness
- After each generation,

$$w_i(t) = w_i(t - 1) + \sum_{j=1}^P \frac{count(i, j) \times fitness_j}{\sum_{k=1}^n count(i, k)}$$



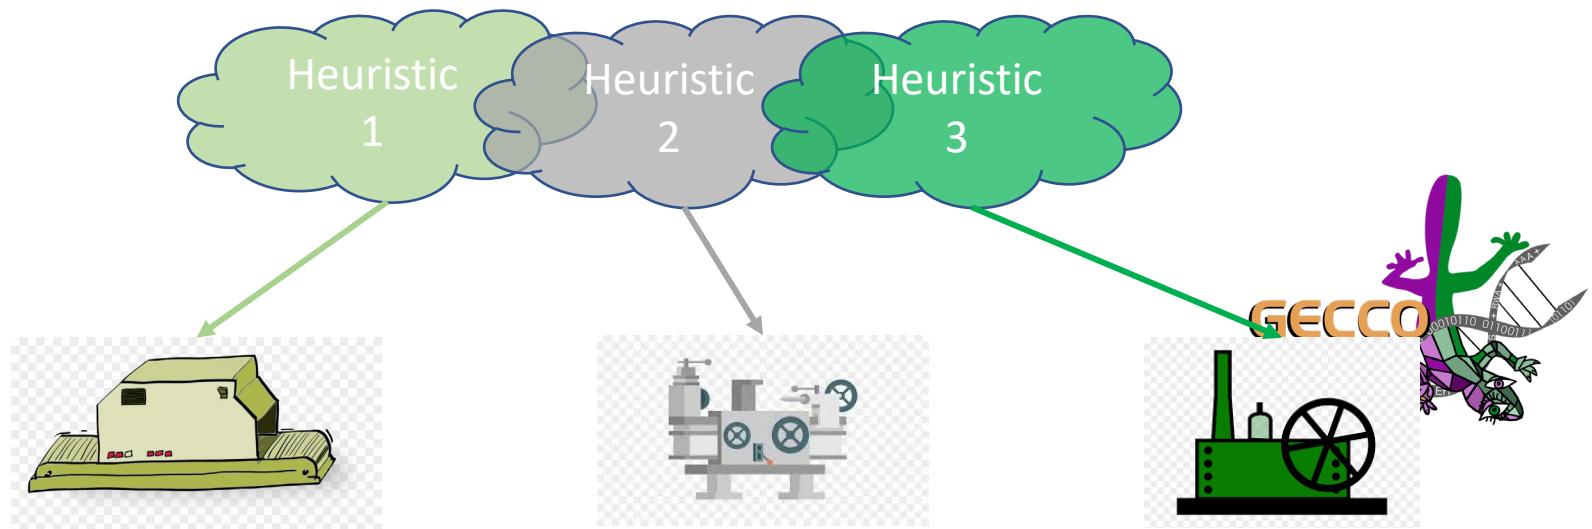
Adaptive Crossover

- Contribution of a sub-tree to the individual based on behaviour consistency (correlation)
- Less contribution will be more likely to be replaced



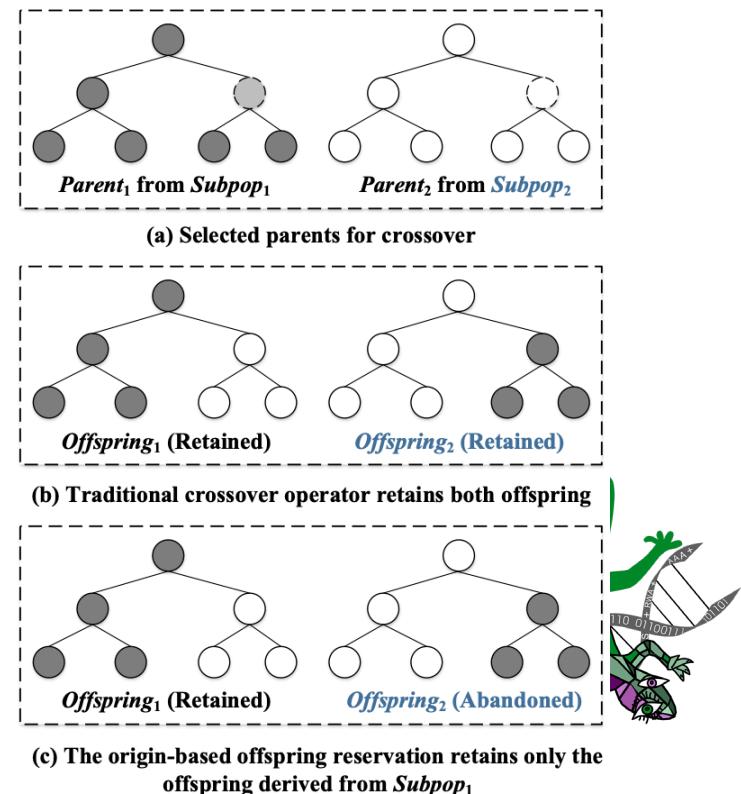
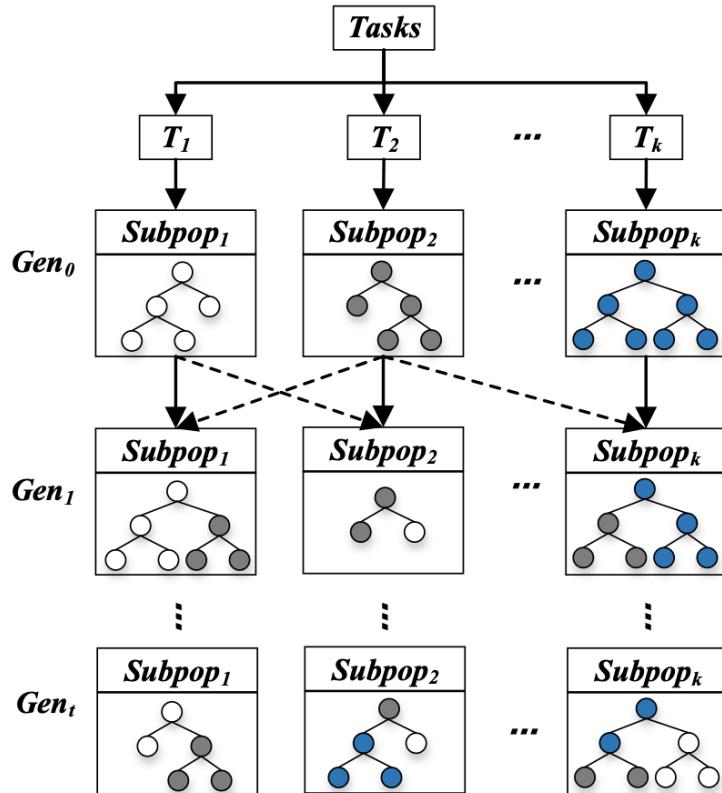
Generality and Beyond

- Hardly possible to have a generalist heuristic for ALL possible scenarios
- Too much to learn specialist heuristic for each scenario
- **Multitask Learning (parallel)**
 - Learn heuristics for multiple scenarios simultaneously
- **Transfer Learning (sequential)**
 - Reuse knowledge to improve learning performance for new scenarios



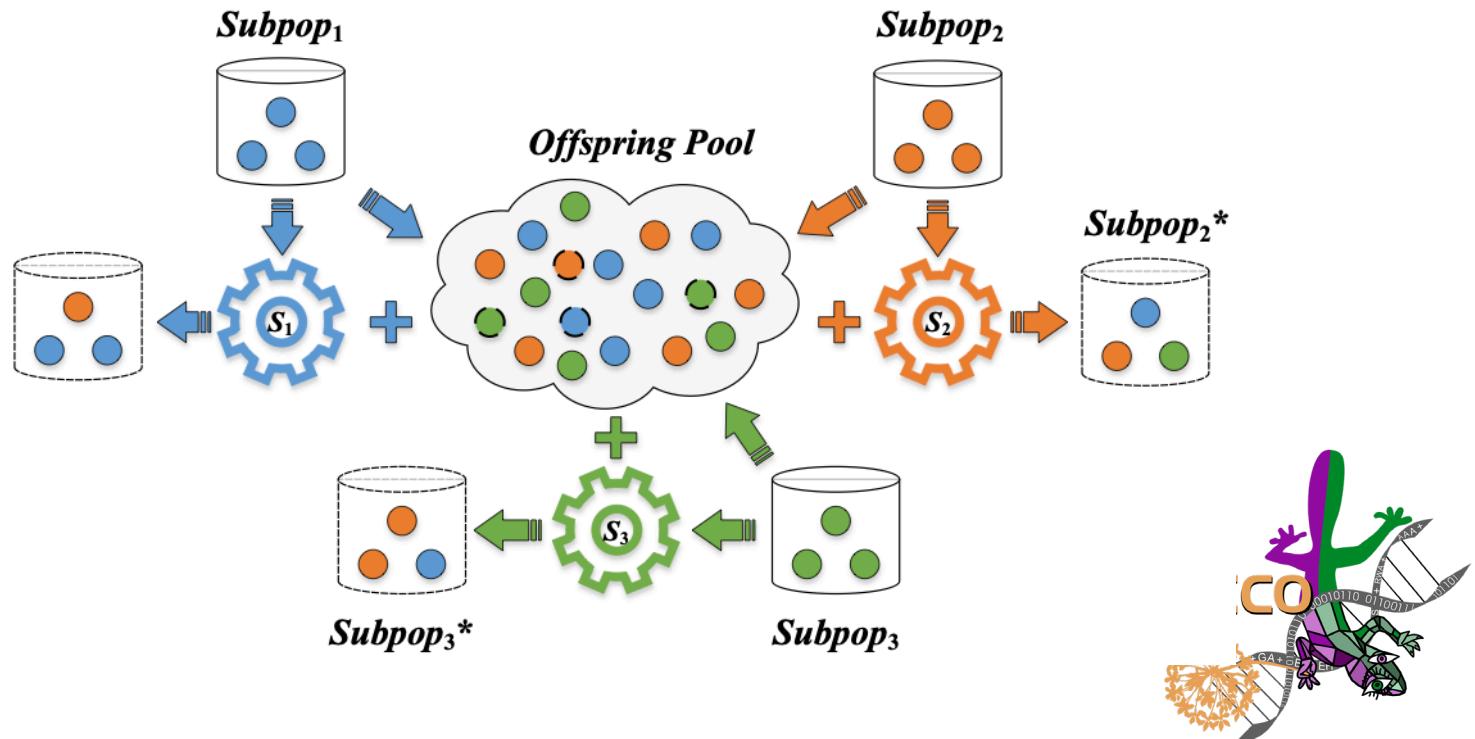
Multitask GP

- Multiple subpopulations, each for a task
- subpopulations interact by crossover, retaining only one offspring more similar to the origin task



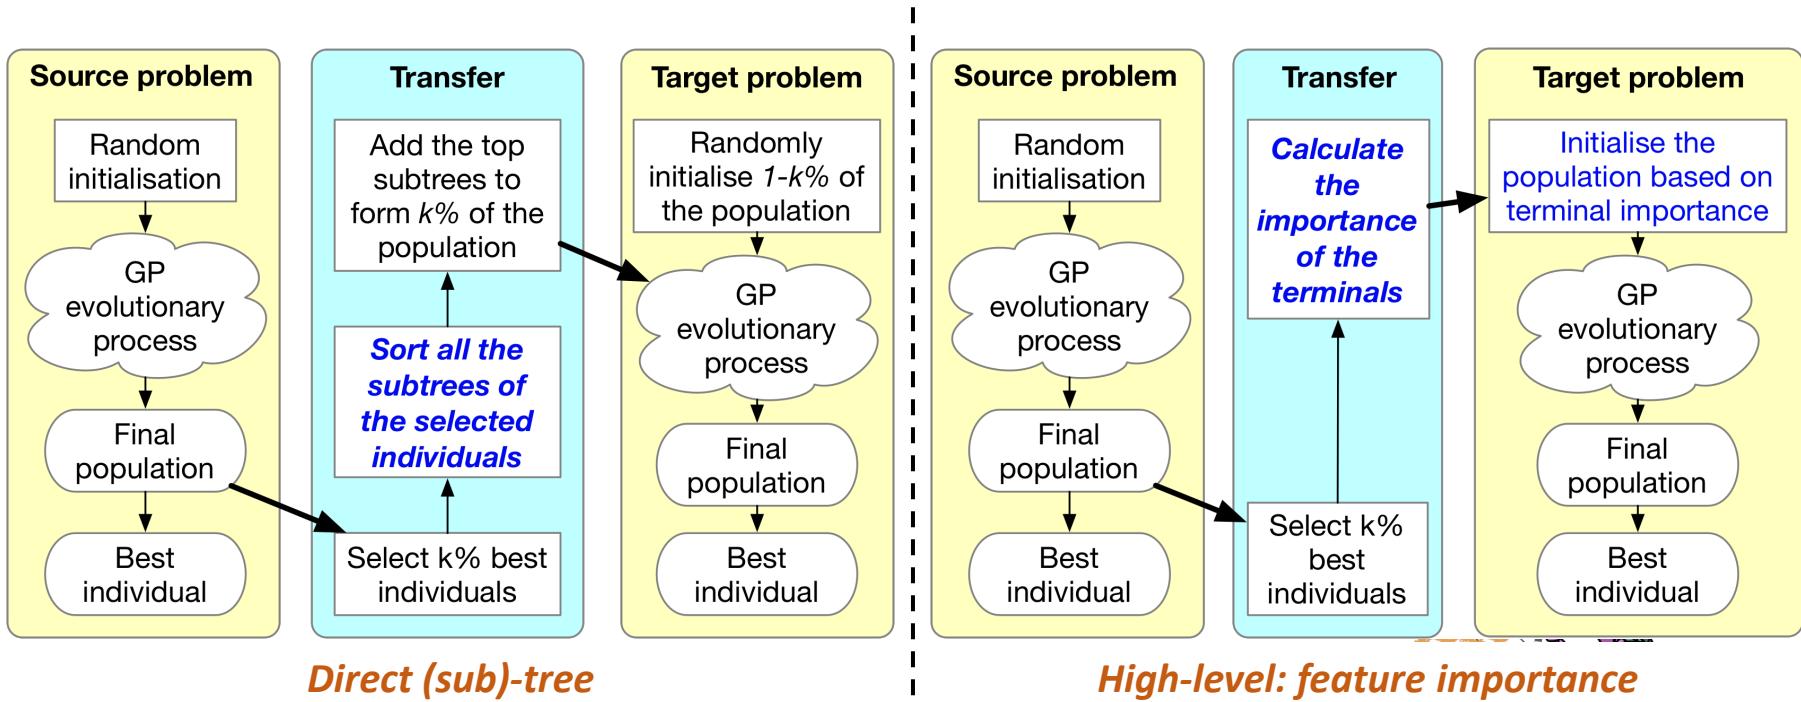
Multitask GP

Migrate offspring among sub-populations based on their effectiveness, evaluated by surrogate



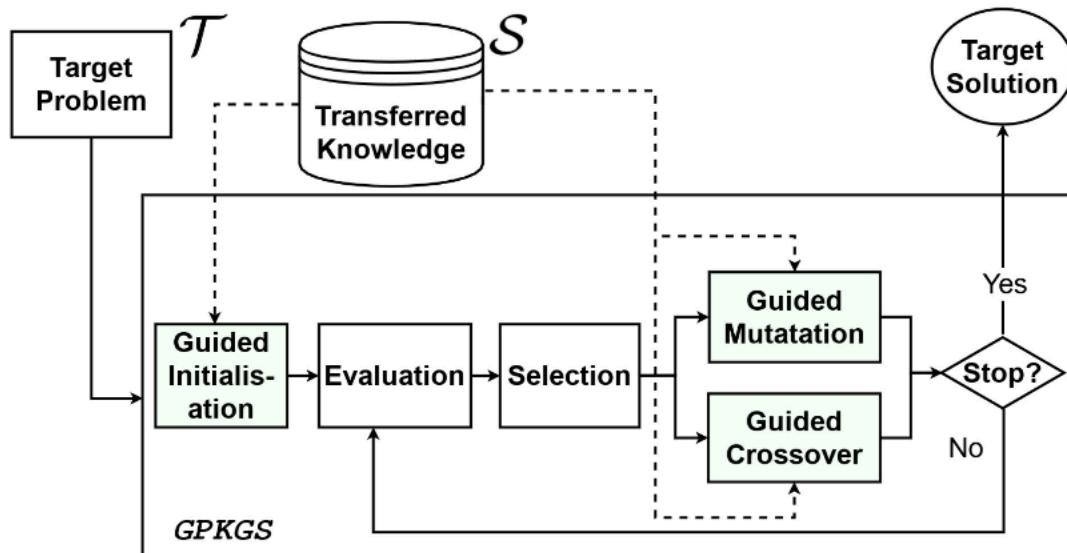
GP with Knowledge Transfer

- What to transfer? Final (sub-)trees, feature importance, ...
- When to transfer? Initialisation, breeding, ...
- How to transfer? Directly transfer, adjust probabilities, ...



GP with Knowledge Transfer

- GP with Knowledge-Guided Search
 - Guided initialization: select the best individuals from the source domain
 - Guided crossover/mutation: to avoid individuals that have been seen in the source domain, to reduce negative transfer



Ardeh, M. A., Mei, Y., & Zhang, M. (2021). Genetic programming with knowledge transfer and guided search for uncertain capacitated arc routing problem. *IEEE Transactions on Evolutionary Computation*, 26(4), 765-779.



Machine Learning to Aid Meta-heuristics (including Evolutionary Methods)

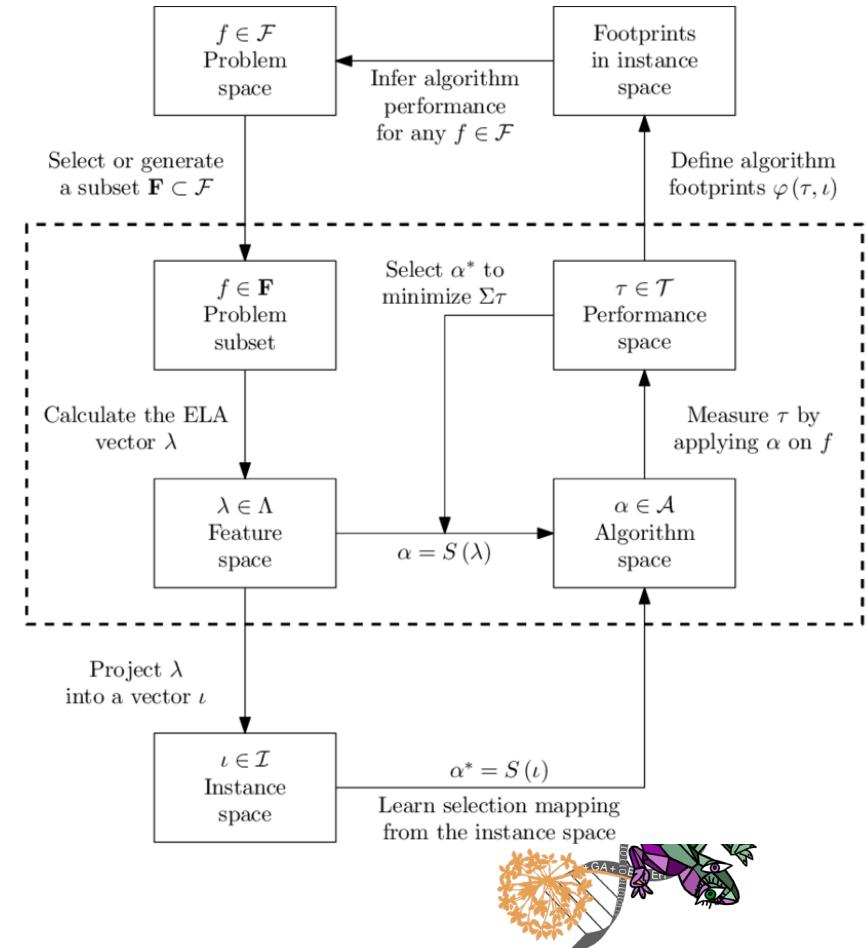
ML to Aid Meta-heuristics

- Learn the **design decisions** in meta-heuristics
 - **Algorithm selection**: how to select the best algorithm for solving a given problem?
 - **Algorithm composition**: how to select the best components of the algorithm, such as initialization method, local search operators, genetic operators?
 - **Algorithm configuration**: how to automatically find the best parameter values?
- **Offline** vs **Online**
- **Supervised** vs **Reinforcement** Learning



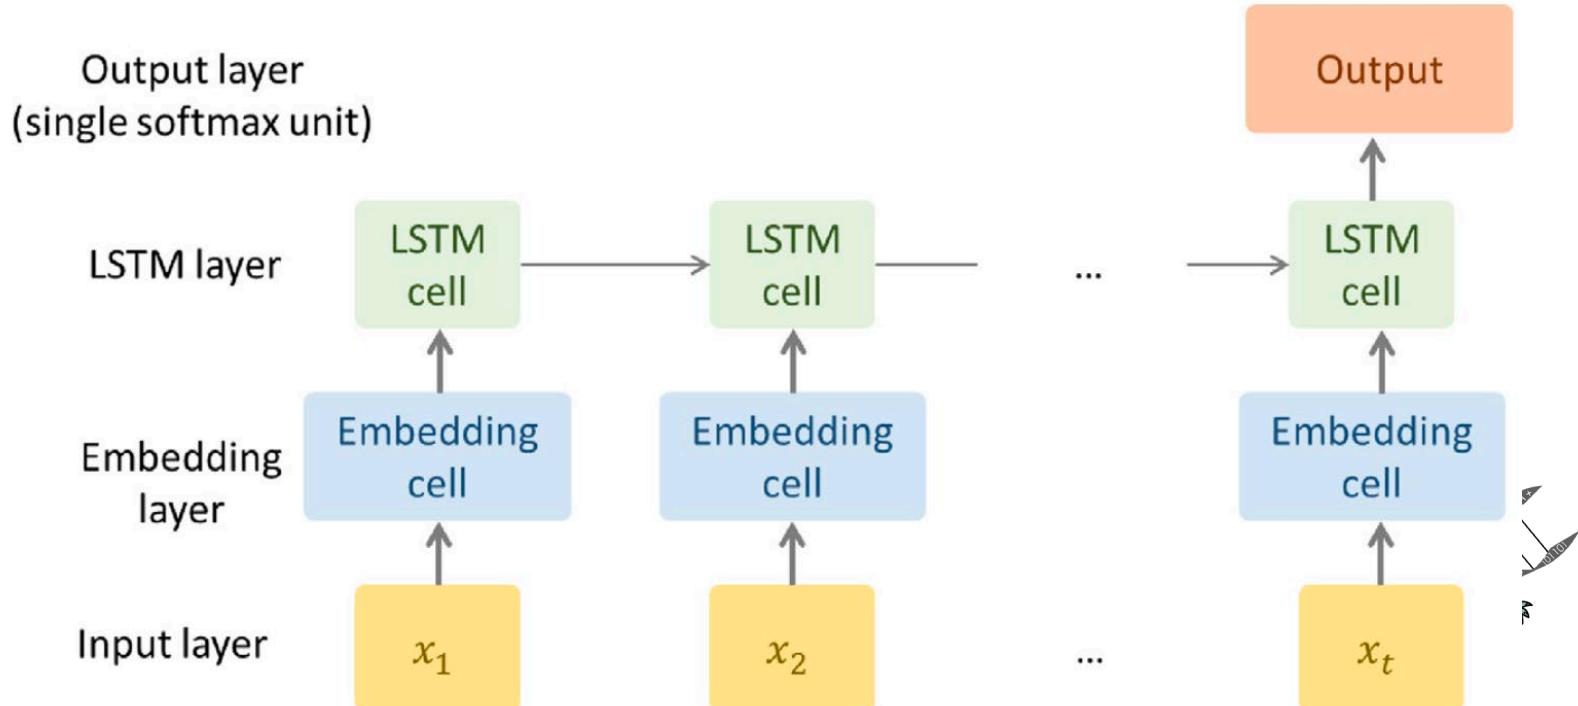
Supervised Offline Algorithm Selection

1. Collect instances and characterise them
2. Collect algorithms
3. Run algorithms to instances
4. Learn model from instance features to selected algorithms, or predicted algorithm performance



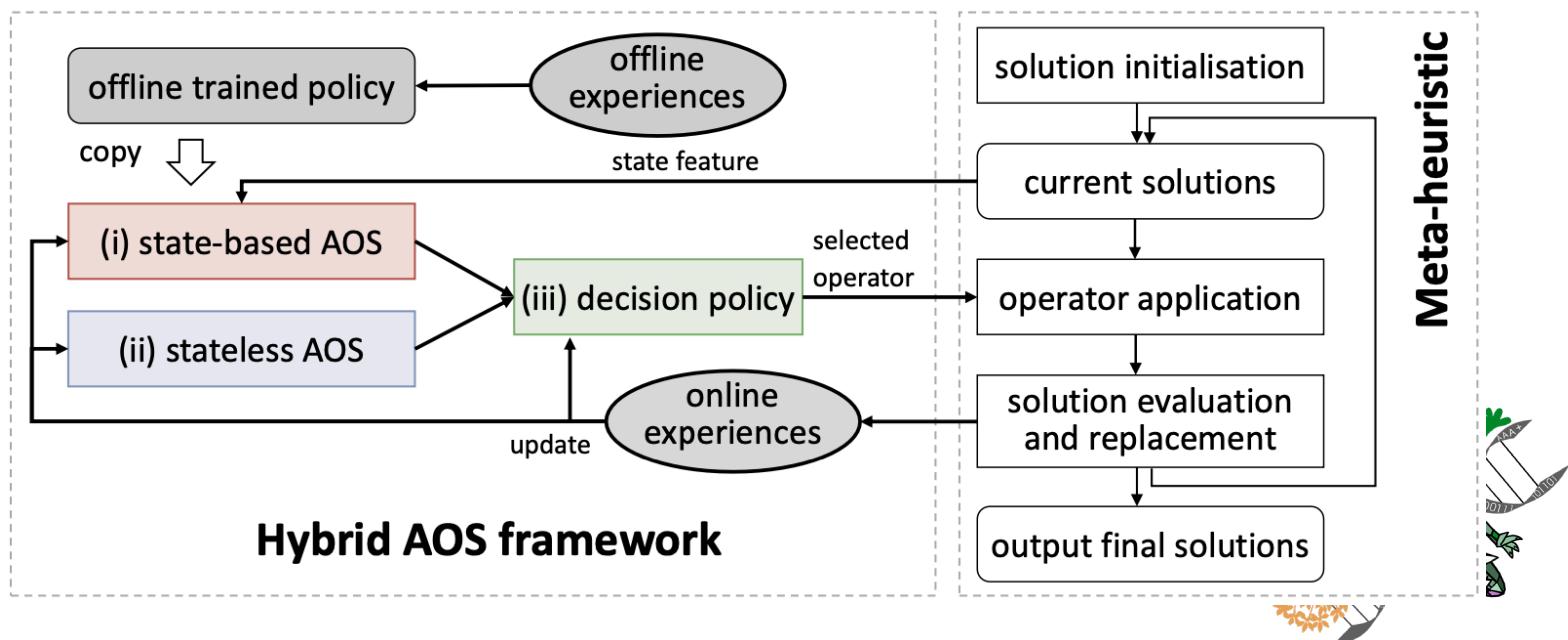
Supervised Offline Algorithm Composition

- Learning operator sequence through LSTM
 - Collect operator sequences
 - Learn LSTM for next operator prediction



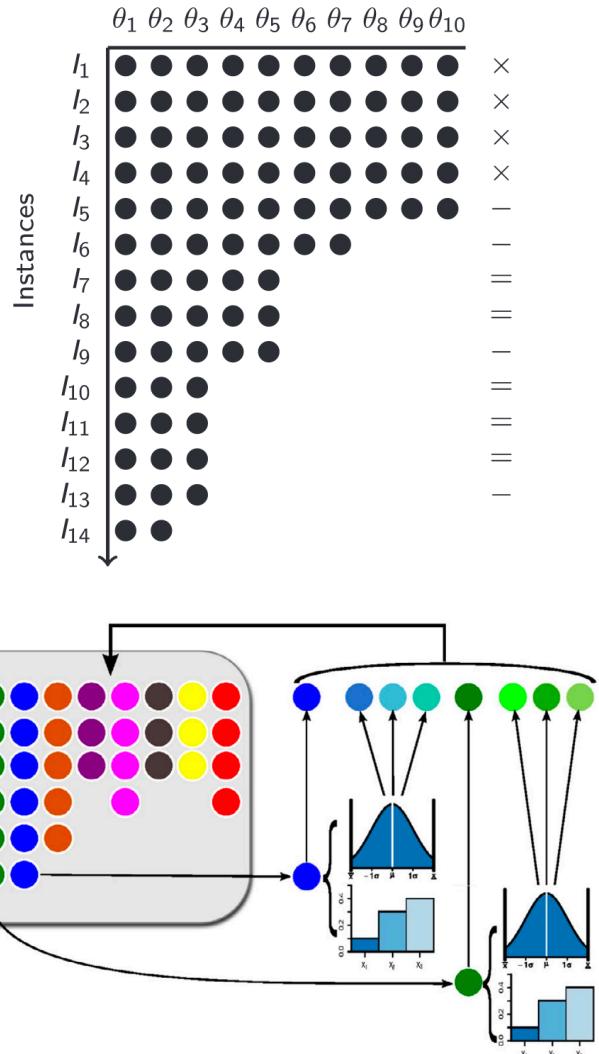
Reinforcement Algorithm Composition

- Offline + Online Adaptive Operator Selection
 - Train **state-based** policy (when to select which operator) offline from training data
 - Update policy online through both **state-based** and **stateless** policy for new problem instance



Reinforcement Algorithm Configuration

- **Irace**: iterative racing for automatic parameter tuning
- **Repeat**:
 - **sampling** new configurations according to a particular distribution,
 - **selecting** the best configurations from the newly sampled ones by means of racing,
 - **updating** the sampling distribution in order to bias the sampling towards the best configurations.



Learning Heatmaps

Learn model indicating **likelihood** for

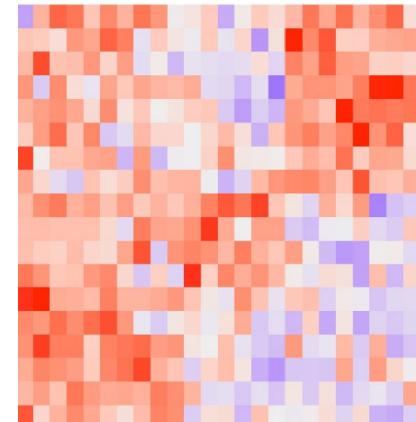
- **components**, e.g., edges in TSP, VRP, nodes in MISp, MaxCut...
- to appear in (close to) optimal solutions

Can be used to

- **prune/sparsen** the search space
- **bias or reduce** operators towards inclusion of more promising components

Training

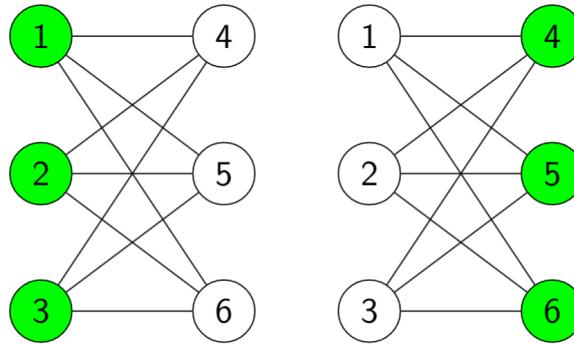
- supervised based on (close to) optimally solved, representative instances
- reinforcement learning



Issues of Heatmaps: 1. Unimodality

Example:

Maximum independent set problem:



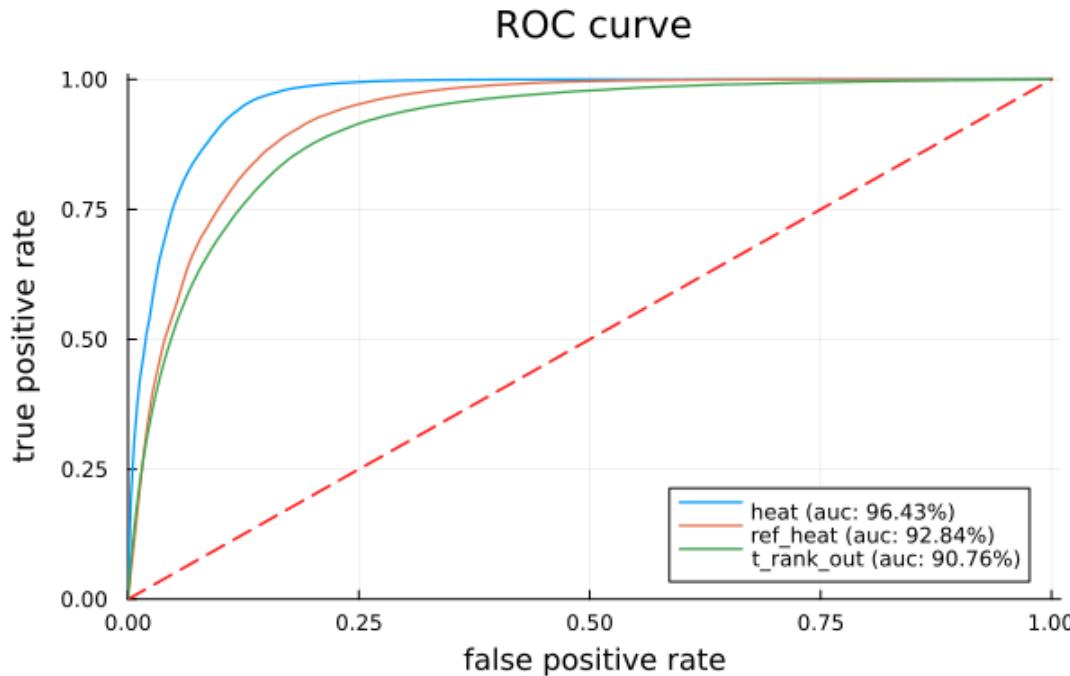
Heatmap: all nodes equally likely in an optimal solution
-> no meaningful information

Symmetries and different (close to) optimal solutions often cause problems.



Issues of Heatmaps:

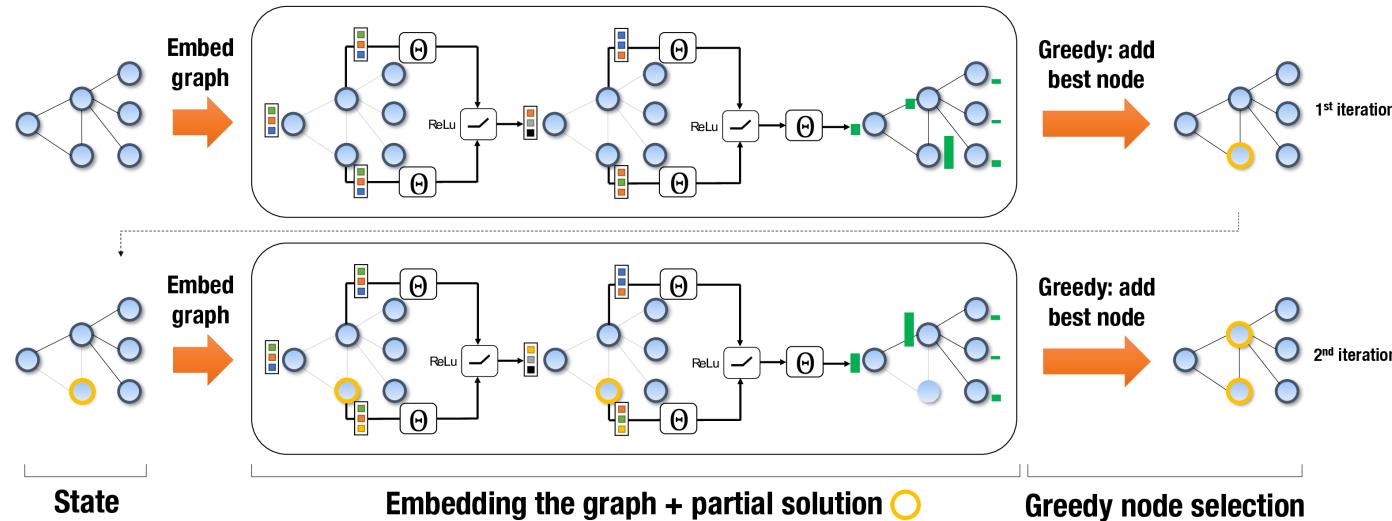
2. Some essential components often do not have characteristics of typical solution components



M. Bresich et al. (2025): Search Space Reduction Through MachineLearning for the Electric Autonomous Dial-A-Ride Problem, submitted



Learning to Solve Graph Problems

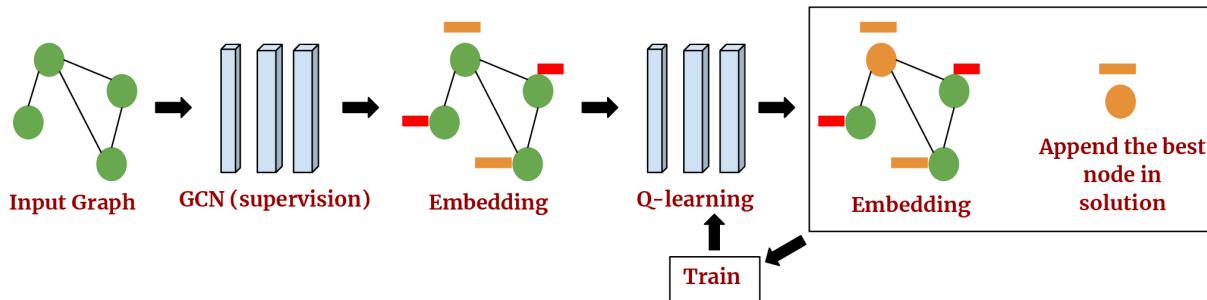


- H. Dai, E.B. Khalil, Y. Zhang, B. Dilkina, and L. Song (2017): [Learning combinatorial optimization algorithms over graphs](#). NeurIPS
- Min vertex cover, max cut, TSP considered
- [Graph neural network \(GNN\)](#) used, structure2vec used to “featurize” nodes
- variant of [Q-learning](#) used to obtain a policy for constructing solutions in an [auto-regressive](#) way



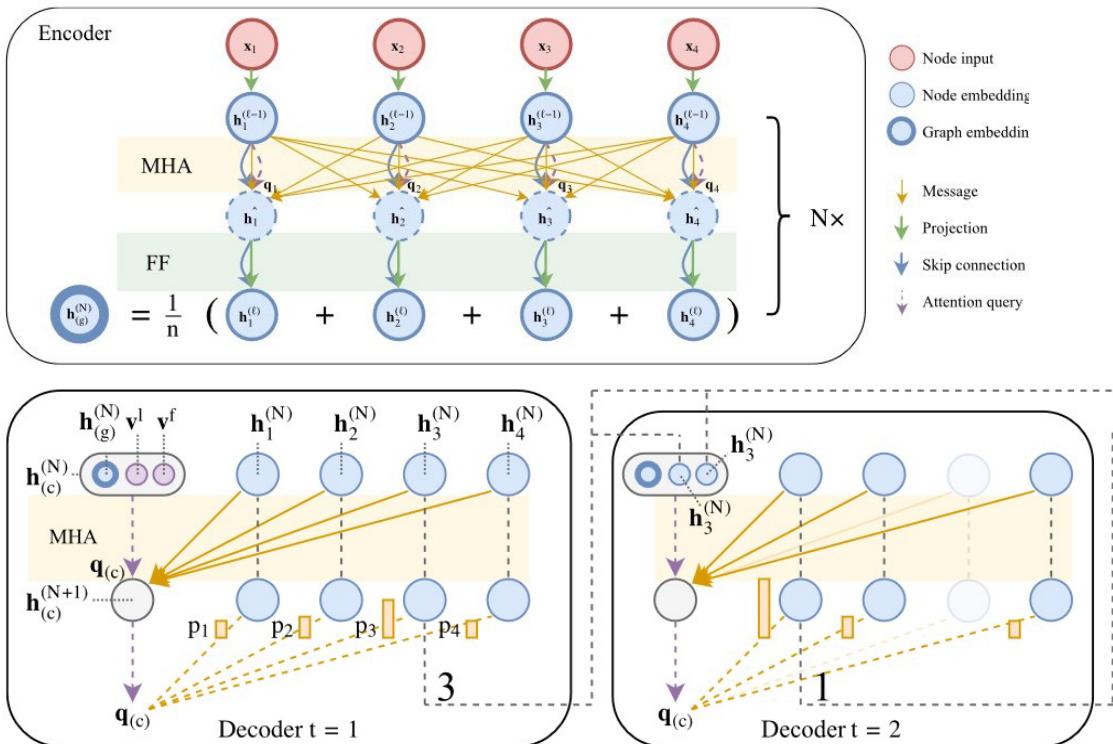
Learning to Solve Graph Problems (cont.)

- A. Mittal, A. Dhawan, S. Manchanda, S. Medya, S. Ranu, A. Singh (2020): [Learning Heuristics over Large Graphs via Deep Reinforcement Learning](#). NeurIPS.
- [focus on scalability](#) to larger graphs,
for max influence, max vertex cover, max coverage problems
- GCN trained in supervised fashion to [predict likelihood](#) of each node to be part of a solution
- Q-learning finds combination of most promising nodes
 $Q(S, v)$: long-term reward for adding node v to solution S
approximated by a simpler GNN



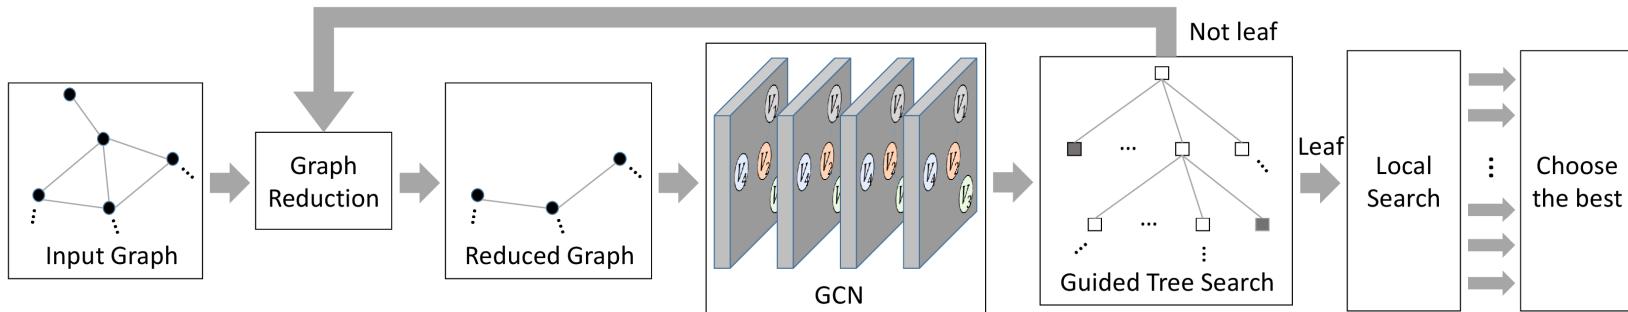
Learning to Solve Graph Problems (cont.)

- W. Kool, H. van Hoof, M. Welling (2019): [Attention, learn to solve routing problems!](#), ICLR
- Autoregressive attention-based encoder + decoder GNN
- for TSP, VRP; trained with REINFORCE



Learning to Solve Graph Problems (cont.)

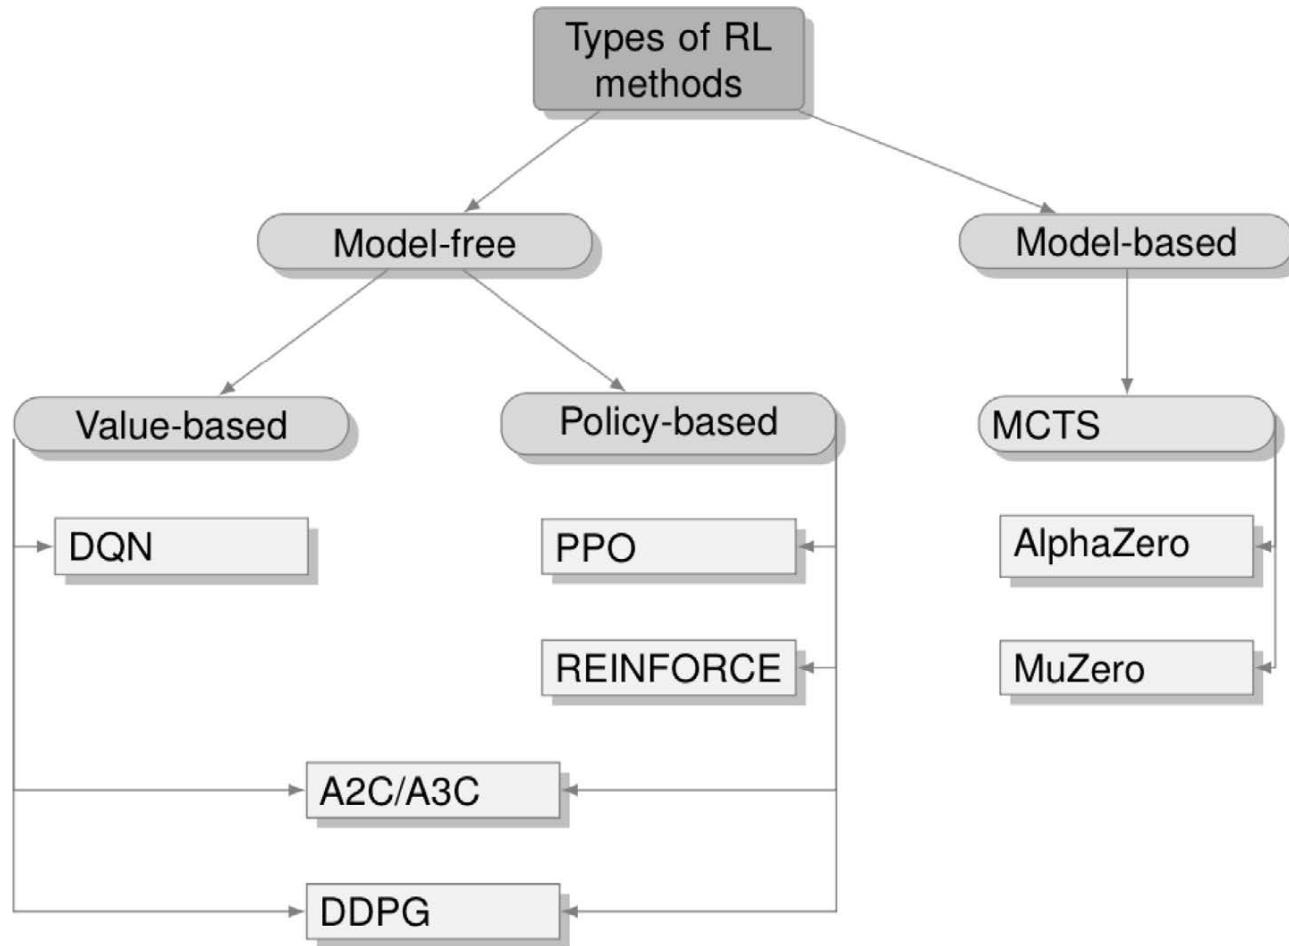
- Z. Li, Q. Chen, V. Koltun (2018): Combinatorial optimization with graph convolutional networks and guided tree search, NeurIPS



- for max independent set, min vertex cover, max clique, SAT
- GNN used to predict likelihood of each node to be part of a solution; yields multiple probability
- heuristic tree search utilizing multiple maps, graph reduction, local search applied
- supervised learning



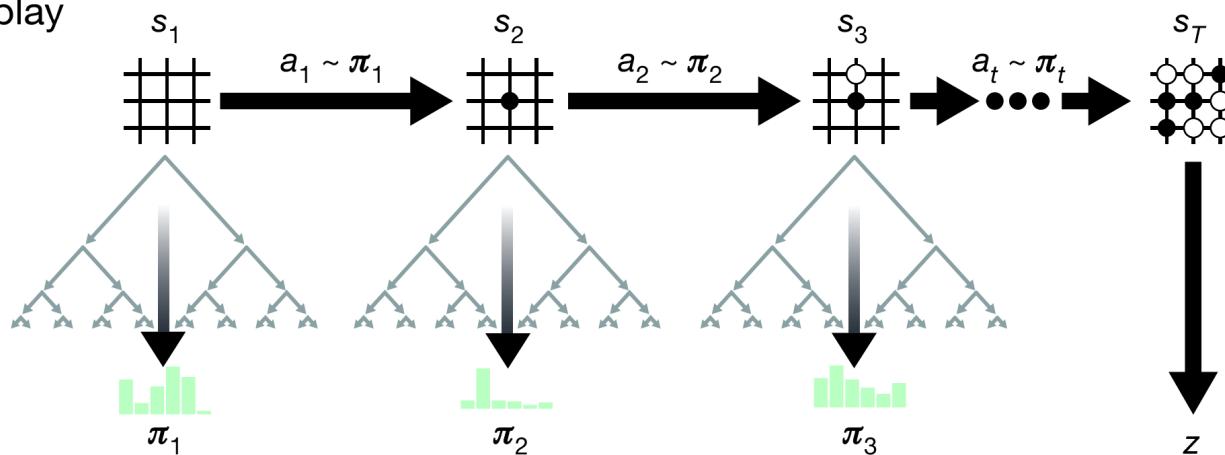
Reinforcement Learning (RL) - Classification



Basic Idea of AlphaZero

- D. Silver et al. (2018). A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 362(6419)
- Learns only by iterated selfplay
- Monte Carlo Tree Search (MCTS) applied to obtain a policy and select a move
- States are evaluated by a Deep NN
Quality of NN output is amplified by MCTS!

Self-play

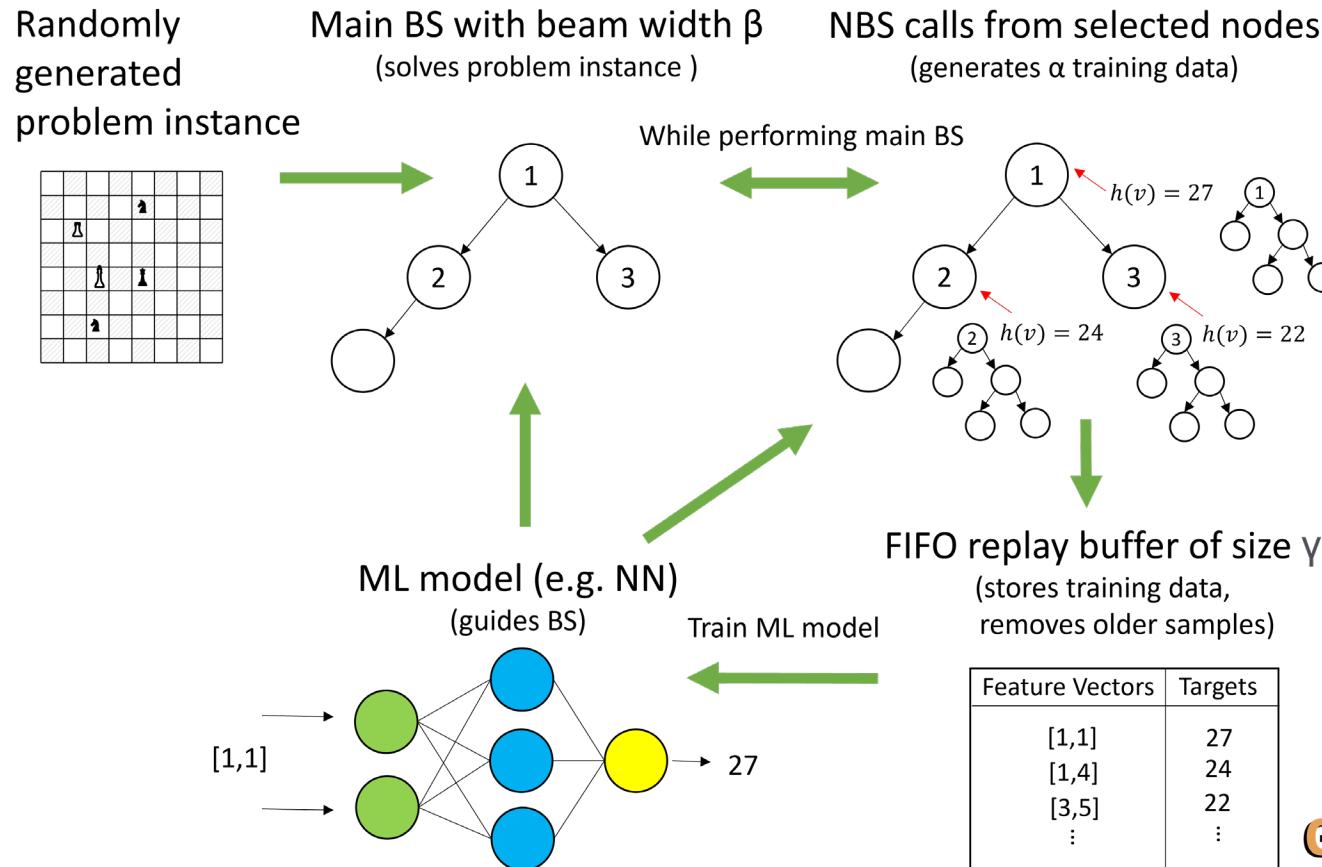


Learning to Solve Graph Problems (cont.)

- K. Abe, Z. Xu, I. Sato, M. Sugiyama, M. (2020): [Solving NP-hard problems on graphs with extended AlphaGo zero](#).arXiv:1905.11623
- for min vertex cover, max cut, max clique problems
- based on AlphaGoZero
- different GNNs tested
- special reward normalization applied
- J. Huang, M. Patwary, and G. Diamos (2019): [Coloring big graphs with AlphaGoZero](#).arXiv:1902.10162
- similar approach for coloring huge graphs
- special FastColorNet neural network architecture



Learning Beam Search (LBS)



M. Huber and G. Raidl (2021). [Learning beam search: Utilizing machine learning to guide beam search for solving combinatorial optimization problems](#). Machine Learning, Optimization, and Data Science (LOD 2021)



LBS for the Longest Common Subsequence Problem

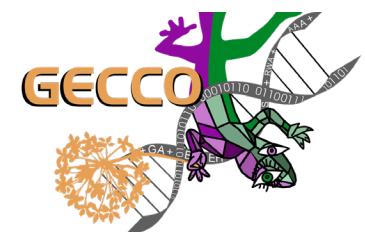
Given: set of m input strings $S = \{s_1, \dots, s_m\}$ over alphabet Σ .

- ▶ Longest Common Subsequence (LCS): find a longest string that appears as subsequence in any string of S .

Example: $m = 2$, $|\Sigma| = 3$

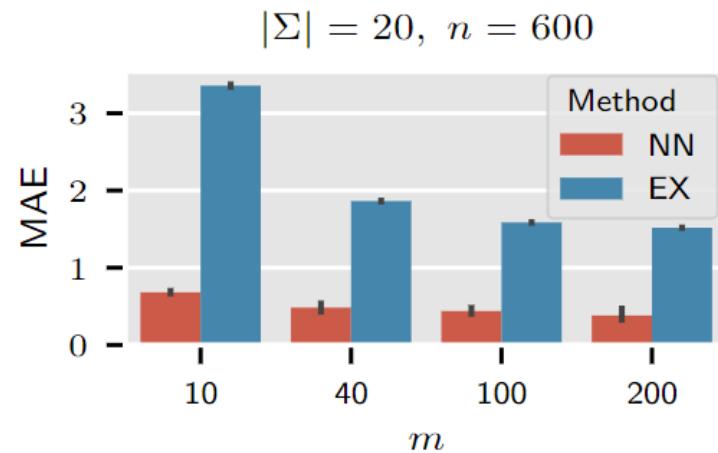
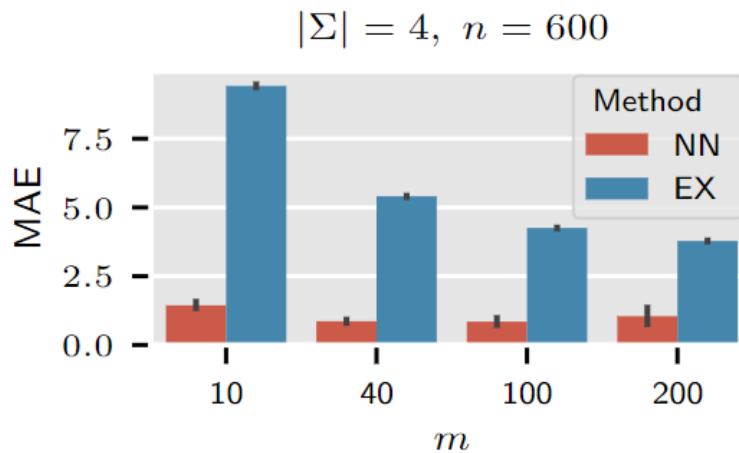
s_1 : ABBA \Rightarrow ABA.
 s_2 : CABA

State-of-the-art: BS with theoretically derived guidance functions EX
(Djukanovic et al., 2020)



LBS Experiments: Approximation of Real LCS Length

The learned network of LBS approximates the real expected LCS lengths better than EX:



LBS Experiments: Results

- Results on rat and BB LCS benchmark instances:
 - NN: MLP with 20+20 hidden nodes
 - Features: remaining input string lengths, remaining min. letter occurrences
 - Beam width:
 - LBS training done with $\beta = 50$
 - Low computation time tests with $\beta = 50$ on [13 out of 28 instances new best results](#)
 - High quality tests with $\beta = 600$ on [7 out of 28 instances new best results](#)
- Also successfully considered: Constrained LCS, shortest common supersequence problem, no-wait flow shop problem

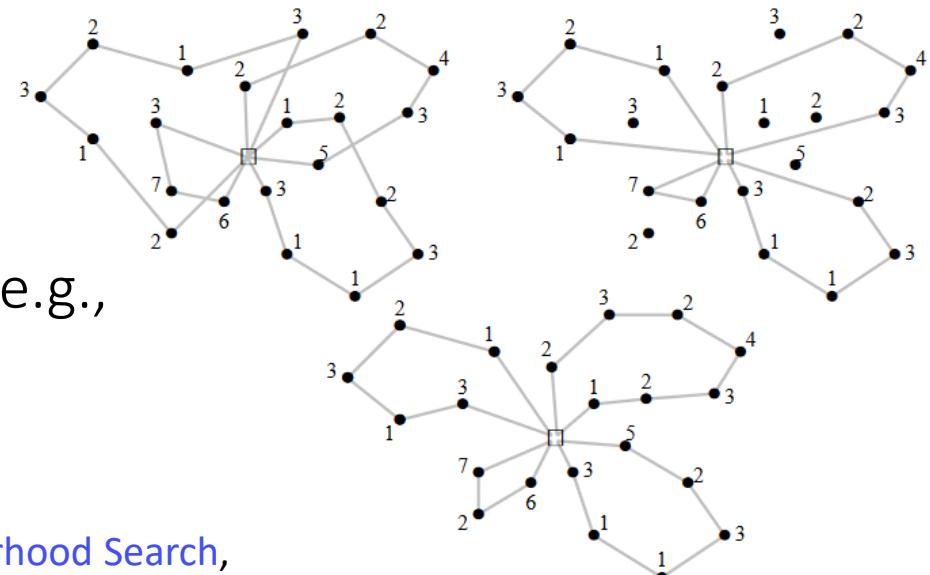


Learning in Large Neighborhood Search

Main principle of Large Neighborhood Search (LNS):

- Iteratively apply **destroy** and **repair** operators
 - Destroy: e.g., remove orders from a VRP tour, remove values assigned to variables
 - Repair: e.g., fast greedy heuristics, dynamic programming, solve a mixed integer linear program

- One of the most prominent metaheuristics, e.g., in transport optimization



D. Pisinger, S. Ropke (2010): Large Neighborhood Search,
Handbook of Metaheuristics, 2nd edition, 399-420, Springer

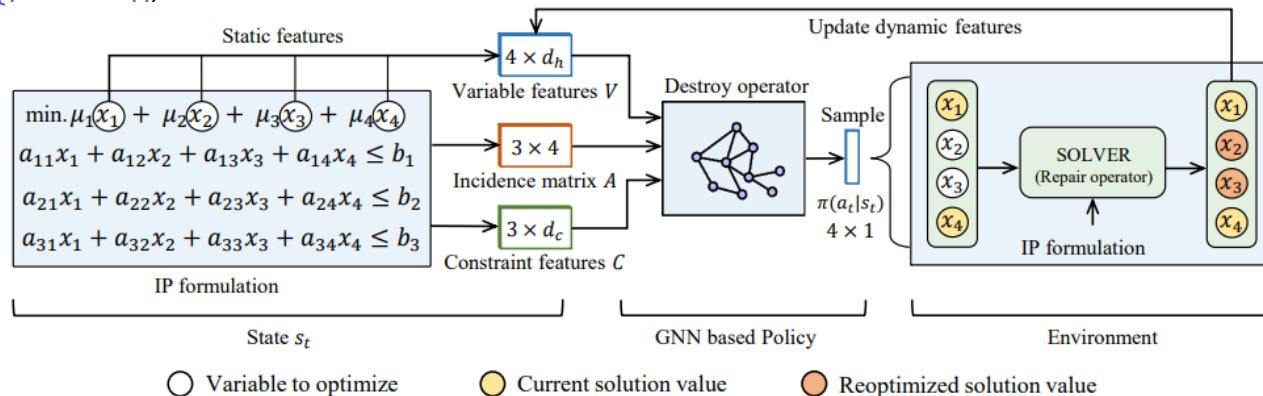
Learning in LNS:

Selection of destroy/repair operators, e.g.,

- J. Kallestad, R. Hasibi, A. Hemmati, K. Sörensen (2023): A [General Deep Reinforcement Learning Hyperheuristic Framework for Solving Combinatorial Optimization Problems](#), EJOR
- S.-N. Johnn, V.-A. Darvariu, J. Handl, J. Kalcsics (2023): [Graph Reinforcement Learning for Operator Selection in the ALNS Metaheuristic](#), OLA.

Selection of what to destroy, e.g.,

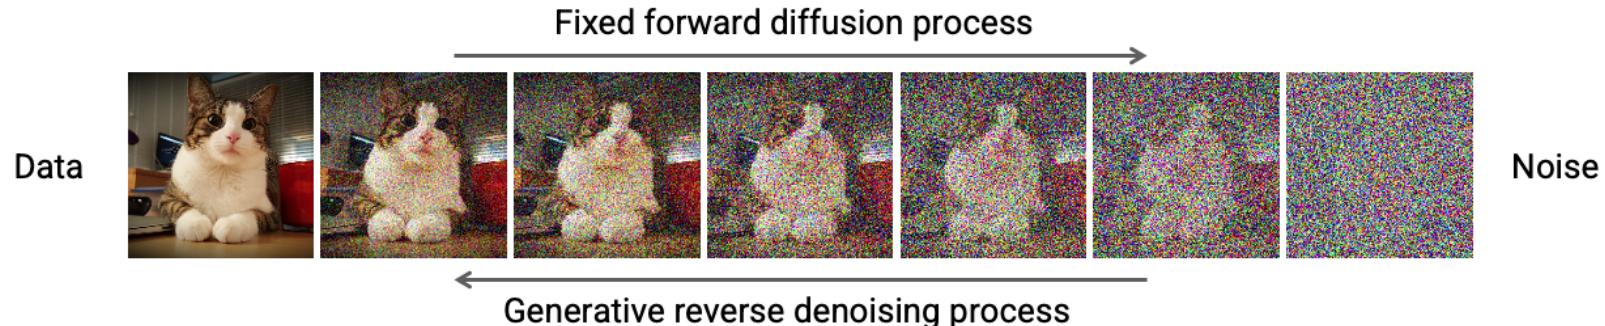
- J. Song, R. Lanka, Y. Yue, B. Dilkina (2020): [A General Large Neighborhood Search Framework for Solving Integer Linear Programs](#), NeurIPS
- R. Addanki, V. Nair, M. Alizadeh (2020): [Neural Large Neighborhood Search](#), NeurIPS
- N. Sonnerat, P. Wang, I. Ktena, S. Bartunov, V. Nair (2021): [Learning a Large Neighborhood Search Algorithm for Mixed Integer Programs](#), arXiv:2107.10201
- Y. Wu, W. Song, Z. Cao, J. Zhang (2021): [Learning Large Neighborhood Search Policy for Integer Programming](#), NeurIPS



- F. Oberweger, G. Raidl, E. Rönnberg, and M. Huber (2022): [A Learning Large Neighborhood Search for the Staff Rerostering Problem](#), CPAIOR
- T. Huang, A. Ferber, Y. Tian, B. Dilkina, B. Steiner (2023): [Searching Large Neighborhoods for Integer Linear Programs with Contrastive Learning](#), ICML

Denoising Diffusion Models (DDMs)

- State-of-the-art in many generative AI applications, in particular the creation of realistically-looking images



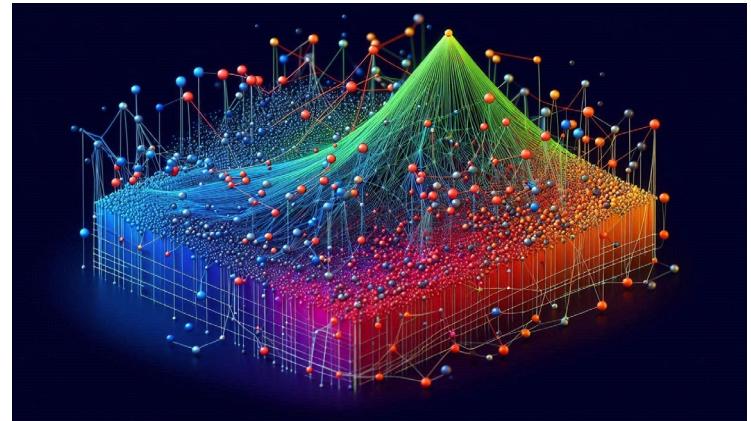
- Training
 - Gaussian noise step-wise added to original images
 - Neural network trained to predict noise added in each step
- Inference
 - Starts from pure random noise
 - Stepwise remove noise via neural network
- DDMs can be conditioned on additional input
- Concept can also be applied to graph neural networks!



DIFUSCO: Graph-Based Diffusion Solver for Comb. Opt.

(Sun and Yang, NeurIPS 2023)

- TSP and MISP considered
- utilizes an anisotropic GNN with edge gating
- discrete diffusion based on Bernoulli noise
- trained on many small instances + (close to) optimal solutions
- (Possible) advantages
 - outperforms earlier approaches by a large margin in their tests
 - faster than autoregressive models
 - better scaling to larger instances
 - multi-modality of solution space is considered



DIFUSCO: Graph-Based Diffusion Solver for Comb. Opt.

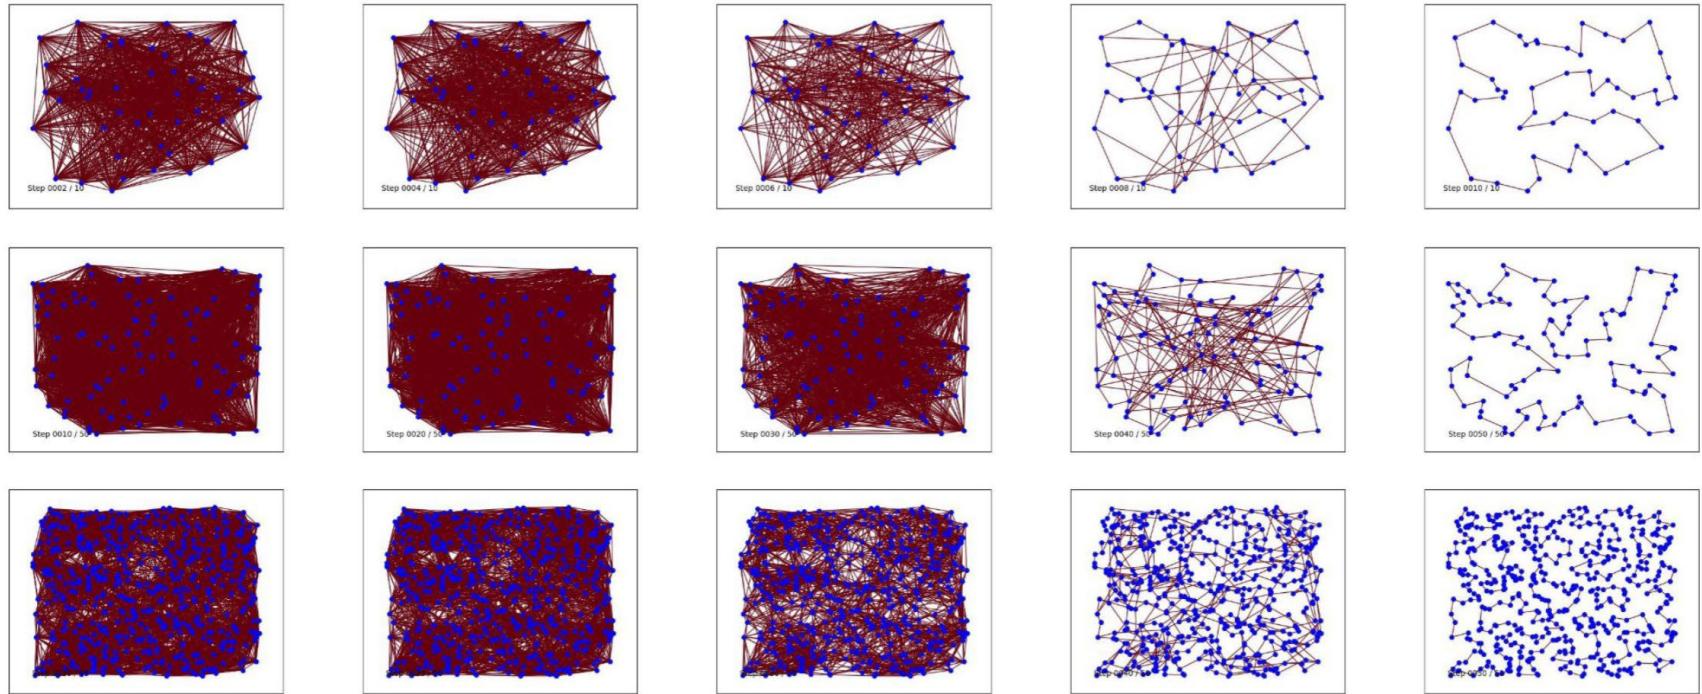


Figure 11: Qualitative illustration of discrete DIFUSCO on TSP-50, TSP-100 and TSP-500 with 50 diffusion steps and cosine schedule.

(from Sun and Yang, 2023)



DIFUSCO: Graph-Based Diffusion Solver for Comb. Opt.

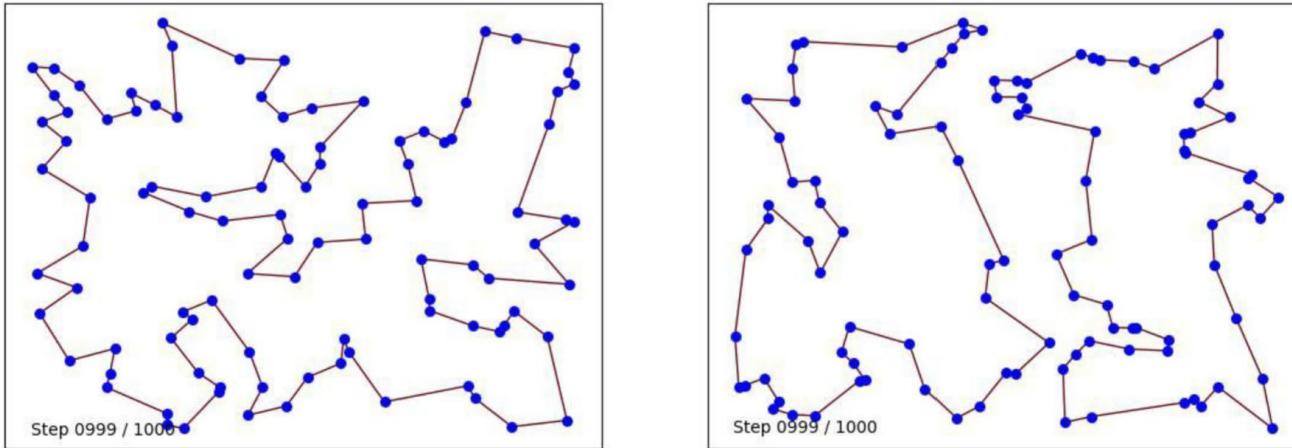
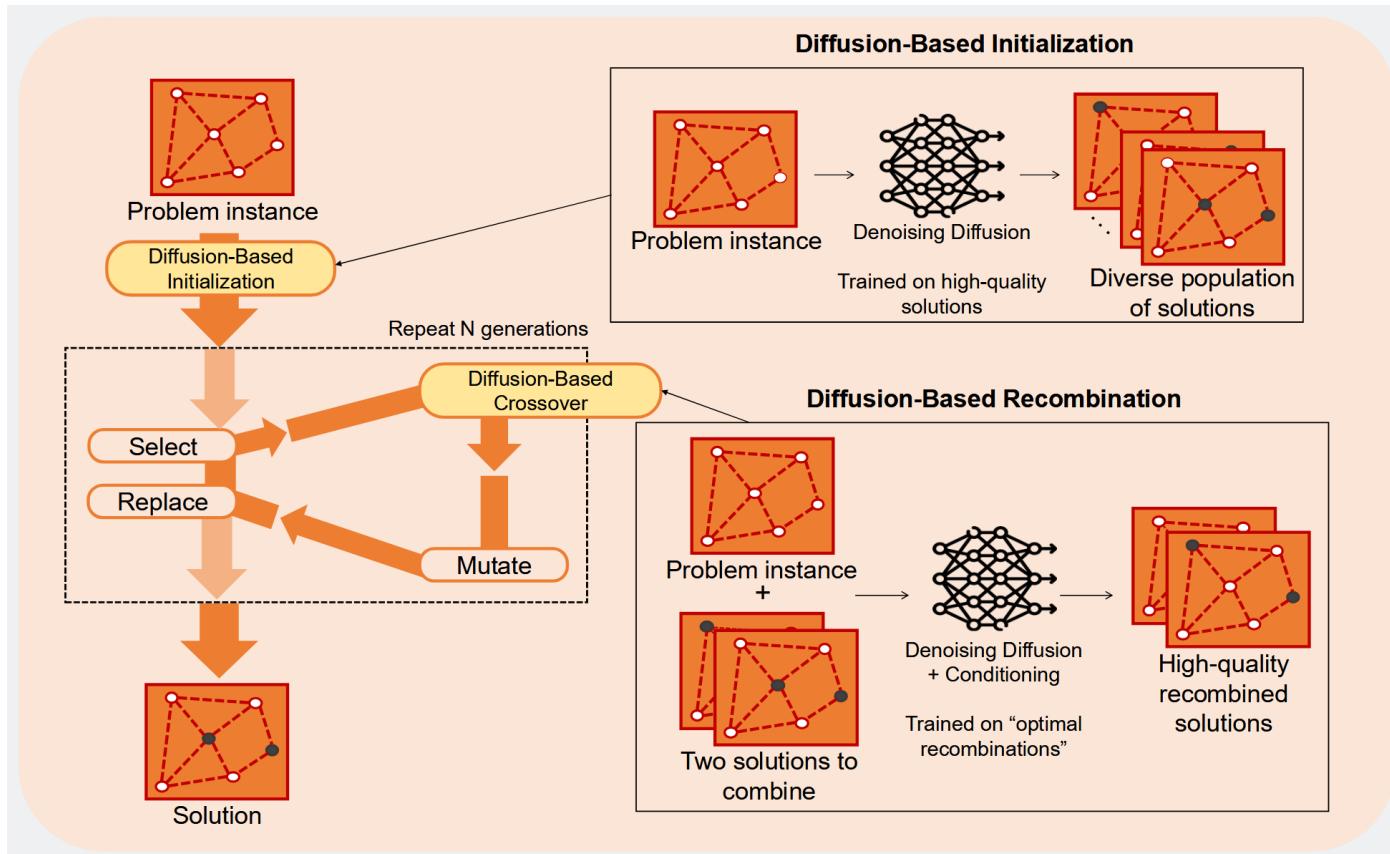


Figure 12: Success (left) and failure (right) examples on TSP-100, where the latter fails to form a single tour that visits each node exactly once. The results are reported without any post-processing.

(from Sun and Yang, 2023)



Denoising Diffusion Based Evolutionary Algorithm



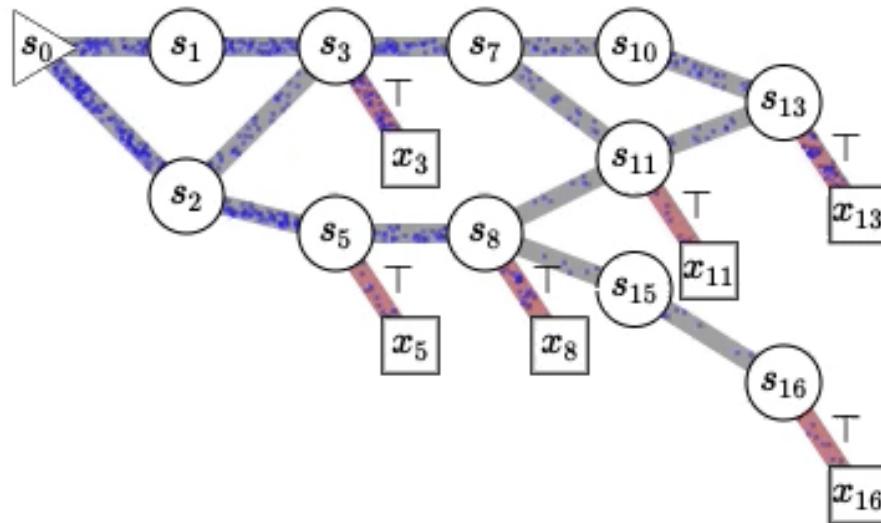
DDEA - Denoising Diffusion Evolutionary Algorithm

J. Salva Soler and G. R. Raidl (2025): A Denoising Diffusion-Based Evolutionary Algorithm Framework Demonstrated on the Maximum Independent Set Problem, LOD 2025

Generative Flow Networks (GFlowNets)

(Bengio, 2021)

- Stochastic policy
- to construct objects through a sequence of steps
- with probability proportional to a reward function



(image from <https://yoshuabengio.org/2022/03/05/generative-flow-networks>)



- Related to, but different from classical RL

Generative Flow Networks (GFlowNets) – contd.

Advantages

- Emphasis on **multi-modality**
- Possibly **more robust training** than classical RL
 - Can be trained off-policy/offline
 - Does not necessarily rely on intermediate rewards
 - Models symmetries more efficiently
- About 60 papers published on GFlowNets since 2021
- **Ongoing work:** utilizing GFlowNets for
 - destroy set generation in LNS
 - as hyper-heuristic to schedule lower-level heuristics





Challenges and Future Directions

Challenges and Future Directions

Highly active research area, learning promising to improve on classical optimization in combinatorial optimization, BUT:

- End-to-end ML approaches are usually not (yet) competitive to carefully crafted classical approaches for COPs
 - Most often, combinations appear most promising
- Generalization to larger instances or instances of other types often is hard
- Obtaining training data and training often very time-consuming
- Deep reinforcement learning does not need labeled data, but often is tricky in practice
- Understand/Interpret the learned algorithms

